

RISC Design: Pipeline Hazards

Virendra Singh

Associate Professor

Computer Architecture and Dependable Systems Lab

Department of Electrical Engineering
Indian Institute of Technology Bombay

<http://www.ee.iitb.ac.in/~viren/>

E-mail: viren@ee.iitb.ac.in

EE-739: Processor Design

Lecture 15 (13 Feb 2013)

CADSL



ILP: Instruction Level Parallelism

- Single-cycle and multi-cycle datapaths execute one instruction at a time.
- How can we get better performance?
- Answer: Execute multiple instruction at a time:
 - **Pipelining** – Enhance a multi-cycle datapath to fetch one instruction every cycle.
 - **Parallelism** – Fetch multiple instructions every cycle.



Traffic Flow



13 Feb 2013

EE-739@IITB

3

CADSL

Pipelining in a Computer

- Divide datapath into nearly **equal tasks**, to be performed serially and requiring non-overlapping resources.
- **Insert registers at task boundaries** in the datapath; registers pass the output data from one task as input data to the next task.
- Synchronize tasks with a clock having a cycle time that just exceeds the time required by the longest task.
- **Break each instruction down into a fixed number** of tasks so that instructions can be executed in a staggered fashion.



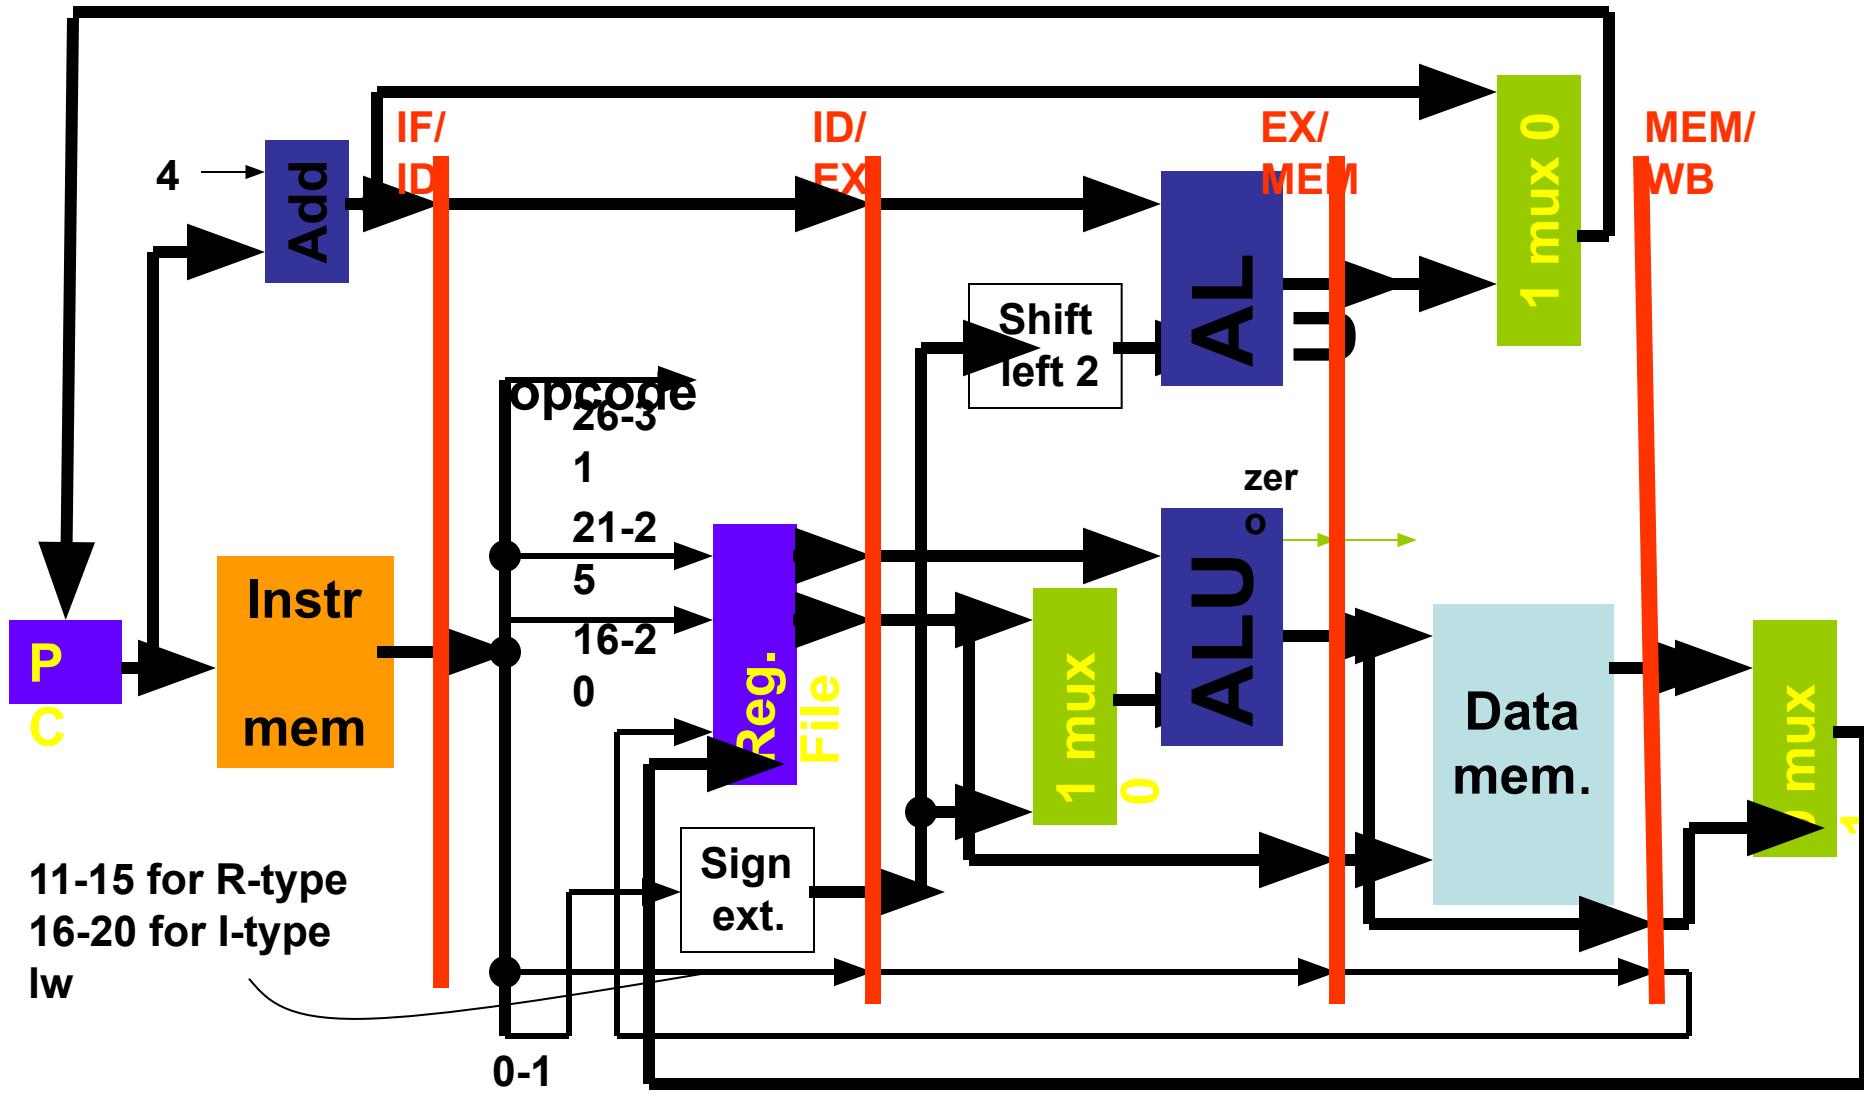
Pipelined Datapath

Instruction class	Instr. fetch (IF)	Instr. Decode (also reg. file read) (ID)	Execution (ALU Operation) (EX)	Data access (MEM)	Write Back (Reg. file write) (WB)	Total time
lw	2ns	1ns 2ns	2ns	2ns	1ns 2ns	10ns
sw	2ns	1ns 2ns	2ns	2ns	1ns 2ns	10ns
R-format: add, sub, and, or, slt	2ns	1ns 2ns	2ns	2ns	1ns 2ns	10ns
B-format: beq	2ns	1ns 2ns	2ns	2ns	1ns 2ns	10ns

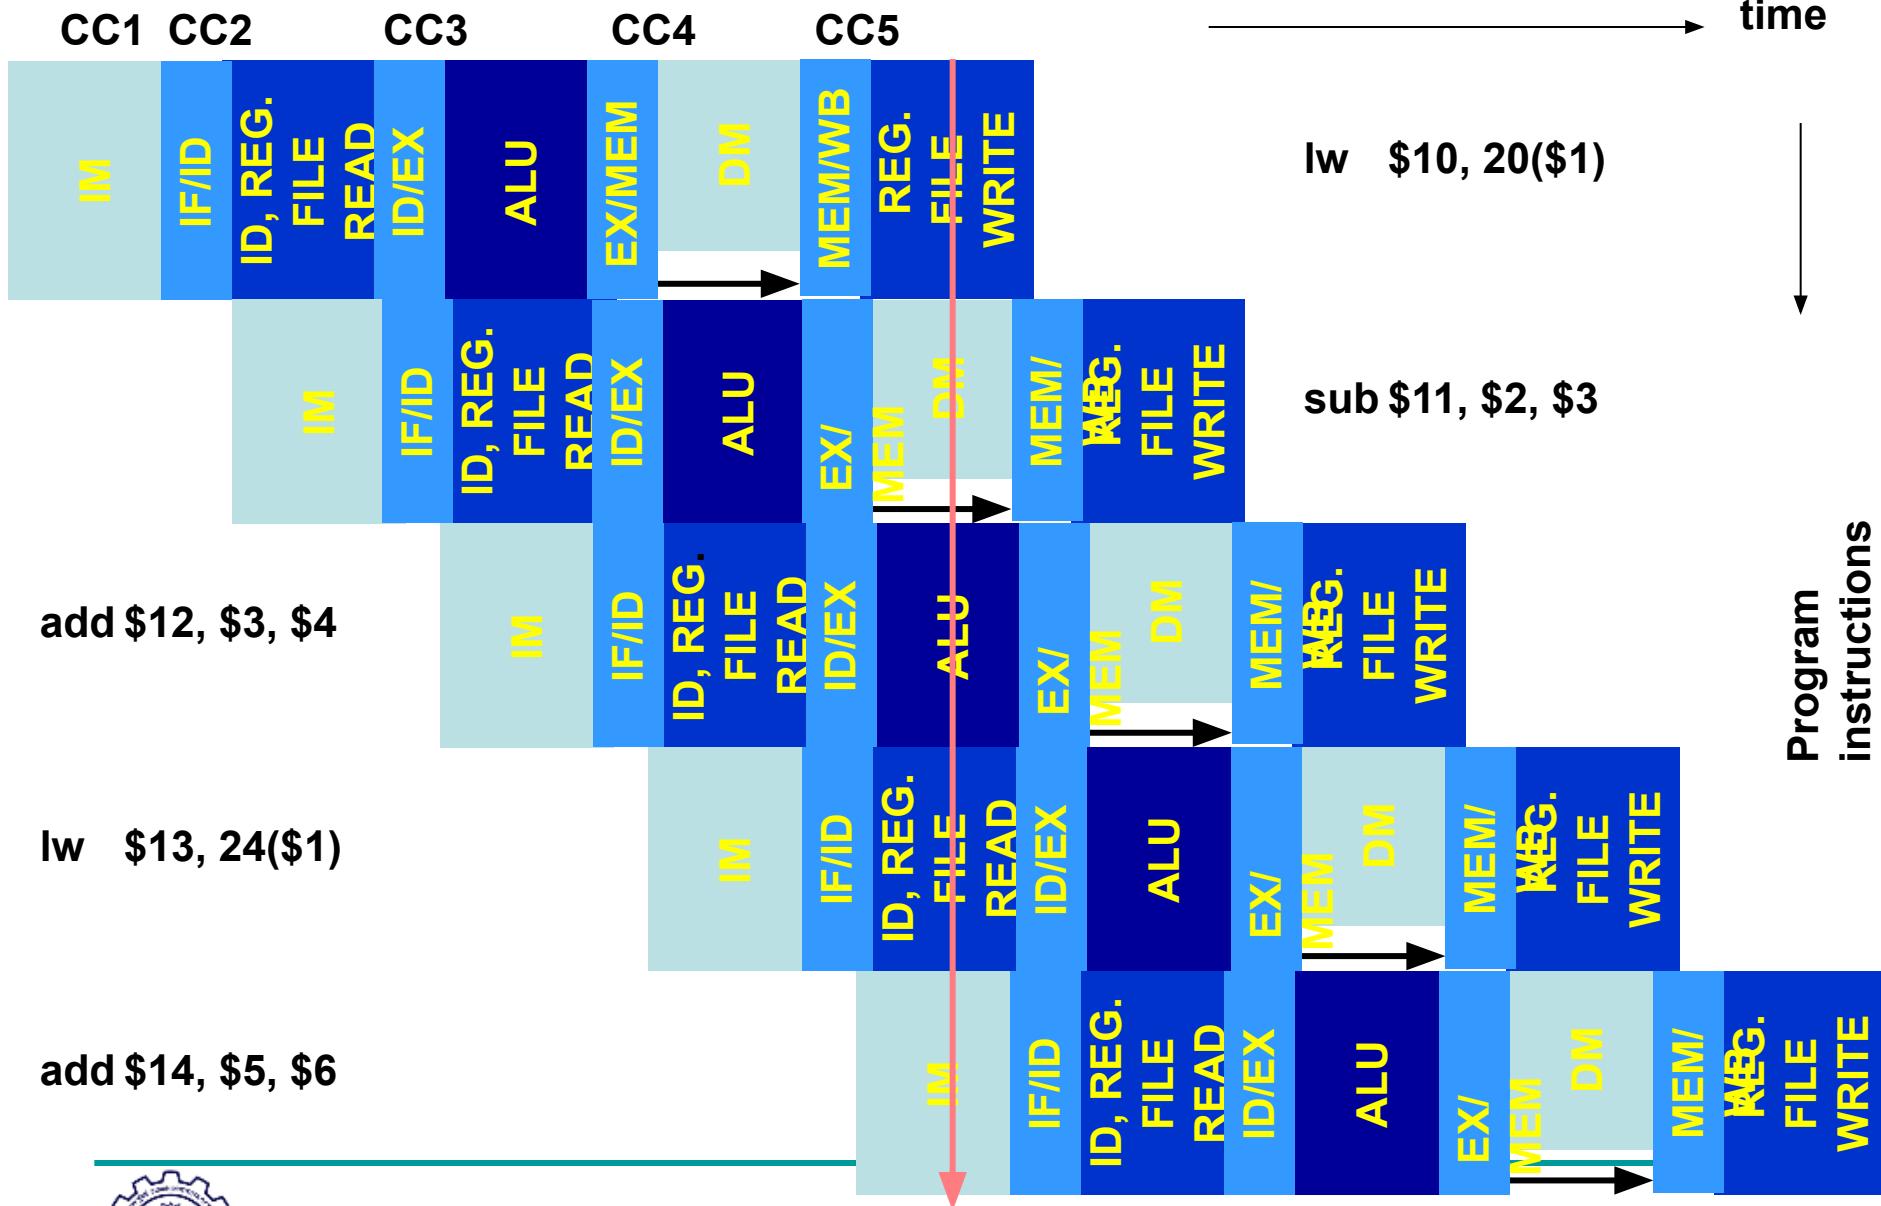
No operation on data; idle time inserted to equalize instruction lengths.



Pipelined Datapath



Program Execution



Single Lane Traffic



Wish list: Highway



CC5

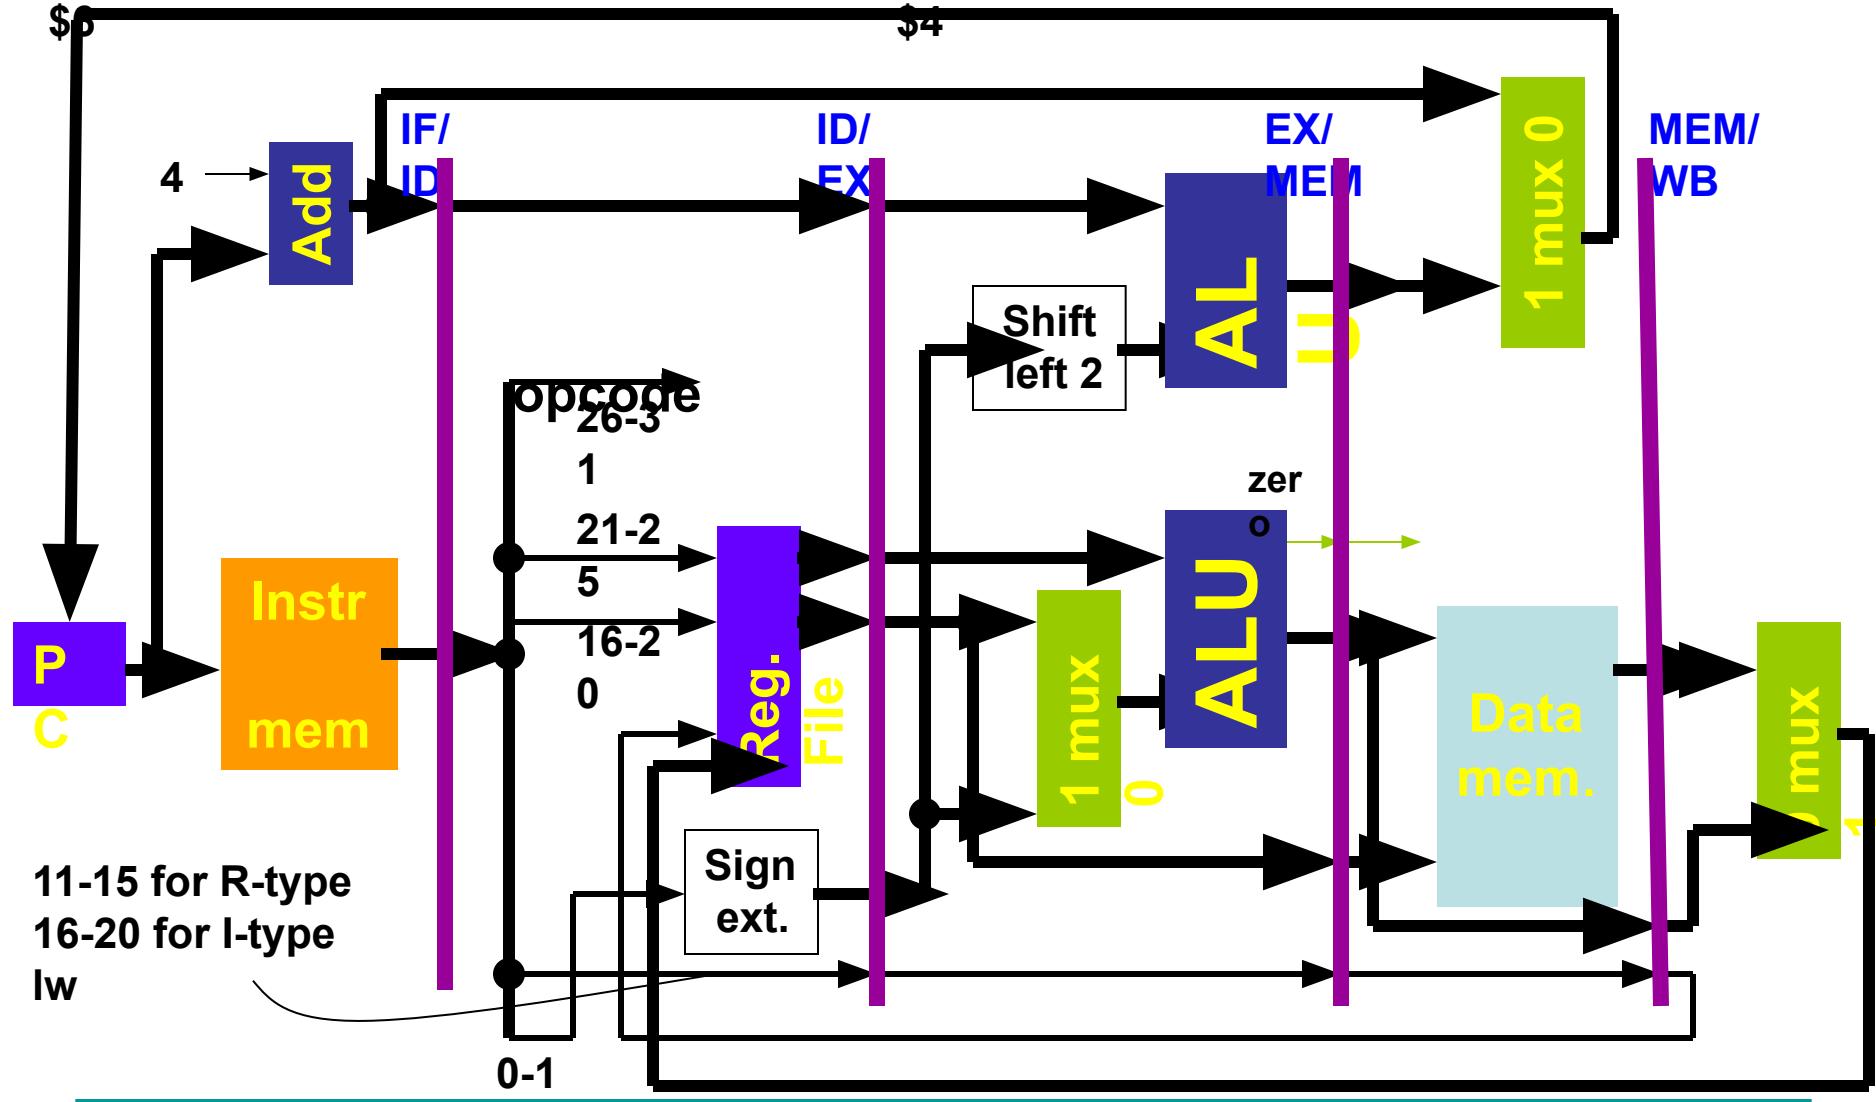
IF: add \$14, \$5,

ID: lw \$13, 24(\$1)

EX: add \$12, \$3,

MEM:
sub \$11, \$2, \$3

WB:
lw \$10, 20(\$1)



Advantages of Pipeline

- After the fifth cycle (CC5), one instruction is completed each cycle; CPI ≈ 1 , neglecting the initial **pipeline latency** of 5 cycles.
 - *Pipeline latency is defined as the number of stages in the pipeline, or*
 - *The number of clock cycles after which the first instruction is completed.*
- The clock cycle time is about four times shorter than that of single-cycle datapath and about the same as that of multicycle datapath.
- For multicycle datapath, CPI = 3.
- So, pipelined execution is faster, but . . .



Science is always wrong. It never solves a problem without creating ten more.

George Bernard Shaw



Pipeline Hazards

- Definition: *Hazard in a pipeline is a situation in which the next instruction cannot complete execution one clock cycle after completion of the present instruction.*
- Three types of hazards:
 - Structural hazard (resource conflict)
 - Data hazard
 - Control hazard



Structural Hazard

- Two instructions cannot execute due to a **resource conflict**.
- Example: Consider a computer with a common data and instruction memory. The fourth cycle of a *lw* instruction requires memory access (memory read) and at the same time the first cycle of the fourth instruction requires instruction fetch (memory read). This will cause a memory resource conflict.



Example of Structural Hazard

CC1 CC2

CC3

CC4

CC5

time

IM/DM

IF/ID

ID, REG.
FILE
READ

ID/EX

ALU

ID, REG.
FILE
READ

ALU

EX/
MEM

IM/DM

MEM/
REG.
FILE
WRITE

lw \$10, 20(\$1)



IM/DM

IF/ID

ID, REG.
FILE
READ

ID/EX

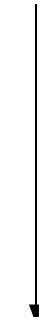
ALU

EX/
MEM

IM/DM

MEM/
REG.
FILE
WRITE

sub \$11, \$2, \$3



Common data
and instr. Mem.
add \$12, \$3, \$4

IM/DM

IF/ID

ID, REG.
FILE
READ

ID/EX

ALU

EX/
MEM

IM/DM

MEM/
REG.
FILE
WRITE

lw \$13, 24(\$1)

IM/DM

IF/ID

ID, REG.
FILE
READ

ID/EX

ALU

EX/
MEM

IM/DM

MEM/
REG.
FILE
WRITE

Needed by two
instructions

Program
instructions

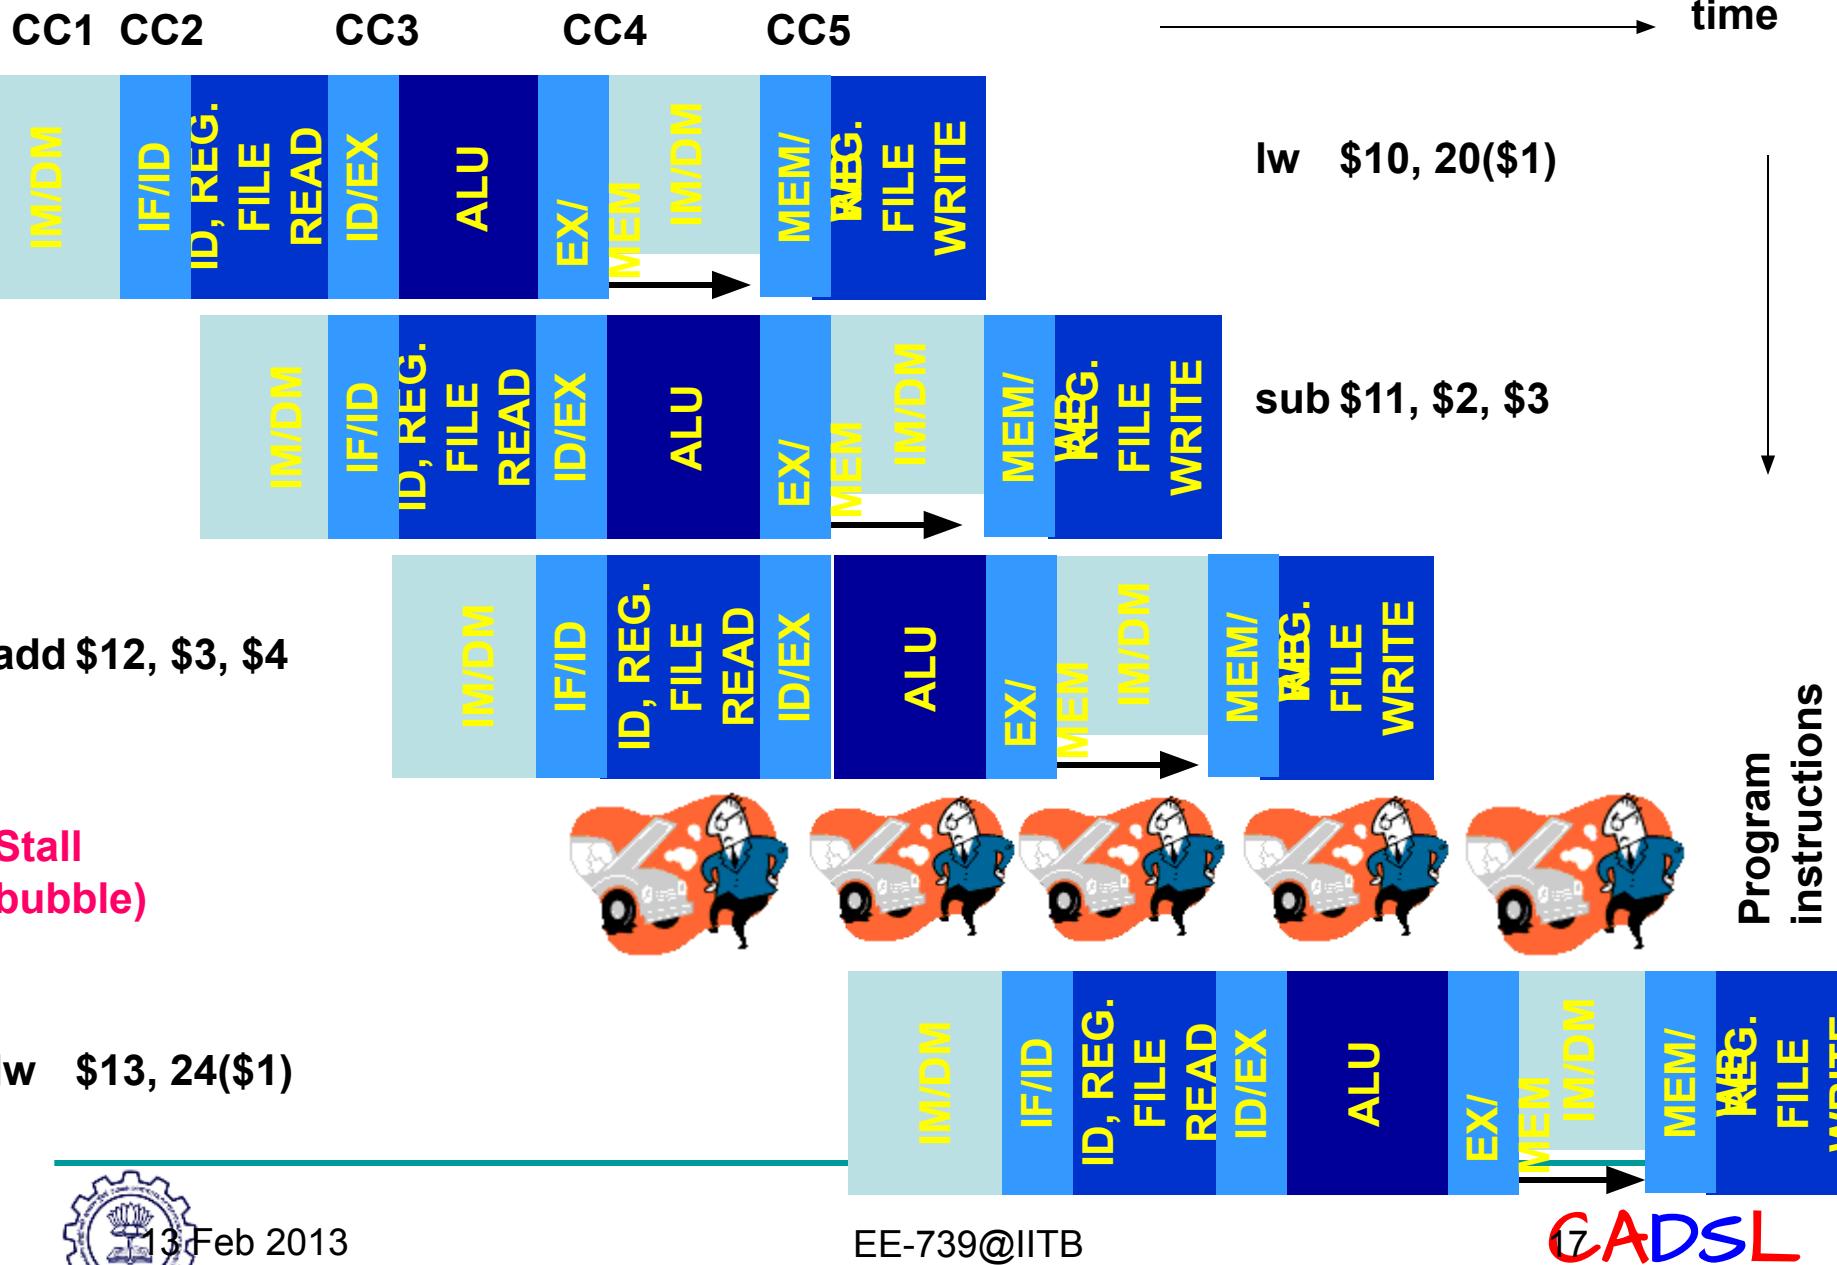


Possible Remedies for Structural Hazards

- Provide duplicate hardware resources in datapath.
- Control unit or compiler can insert delays (no-op cycles) between instructions. This is known as pipeline *stall* or *bubble*.



Stall (Bubble) for Structural Hazard

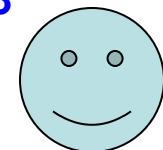
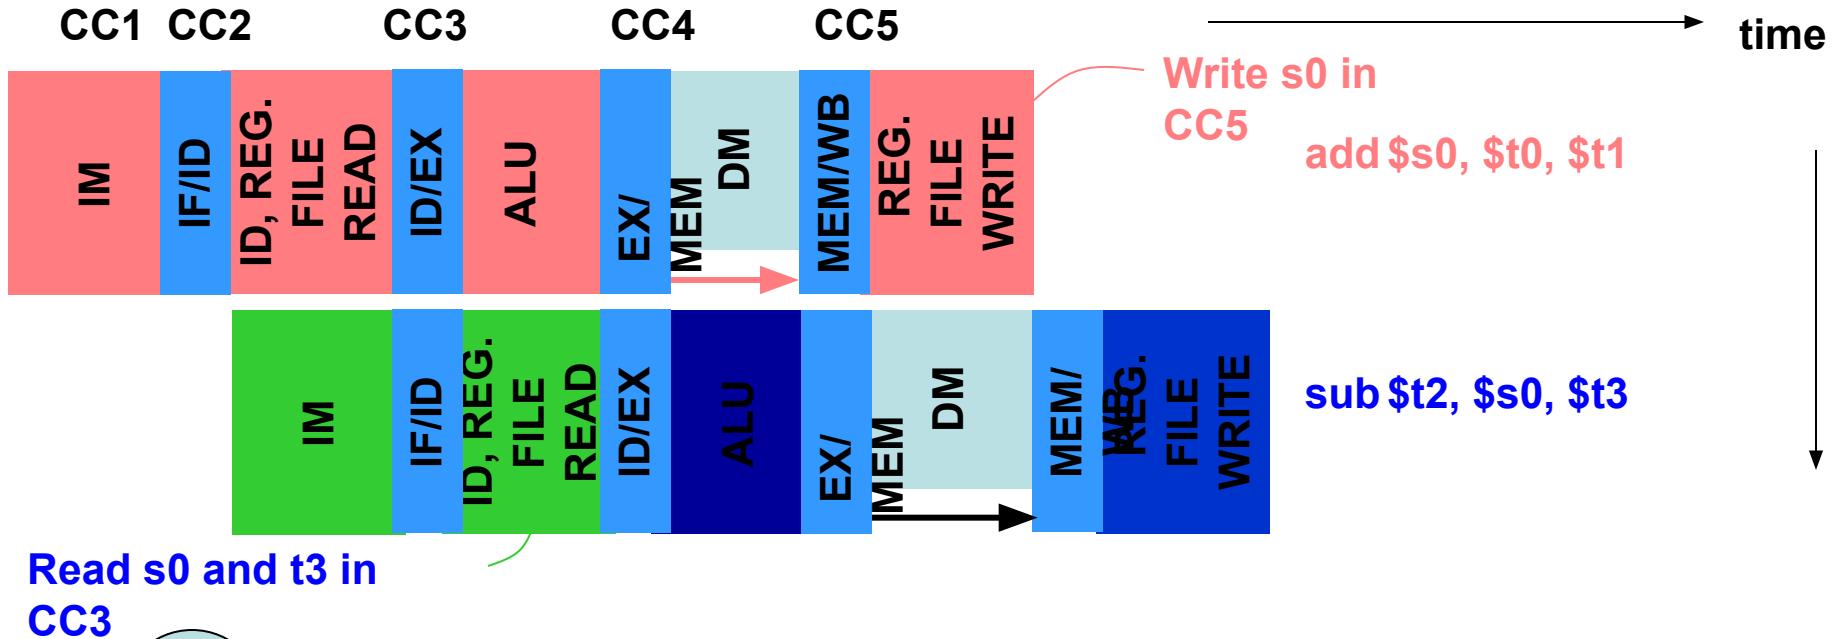


Data Hazard

- Data hazard means that an instruction cannot be completed because the needed data, to be generated by another instruction in the pipeline, is not available.
- Example: consider two instructions:
 - ❖ add \$s0, \$t0, \$t1
 - ❖ sub \$t2, \$s0, \$t3 # needs \$s0



Example of Data Hazard



We need to read s0 from reg file in cycle 3
But s0 will not be written in reg file until cycle 5



However, s0 will only be used in cycle 4
And it is available at the end of cycle 3

Program
instructions

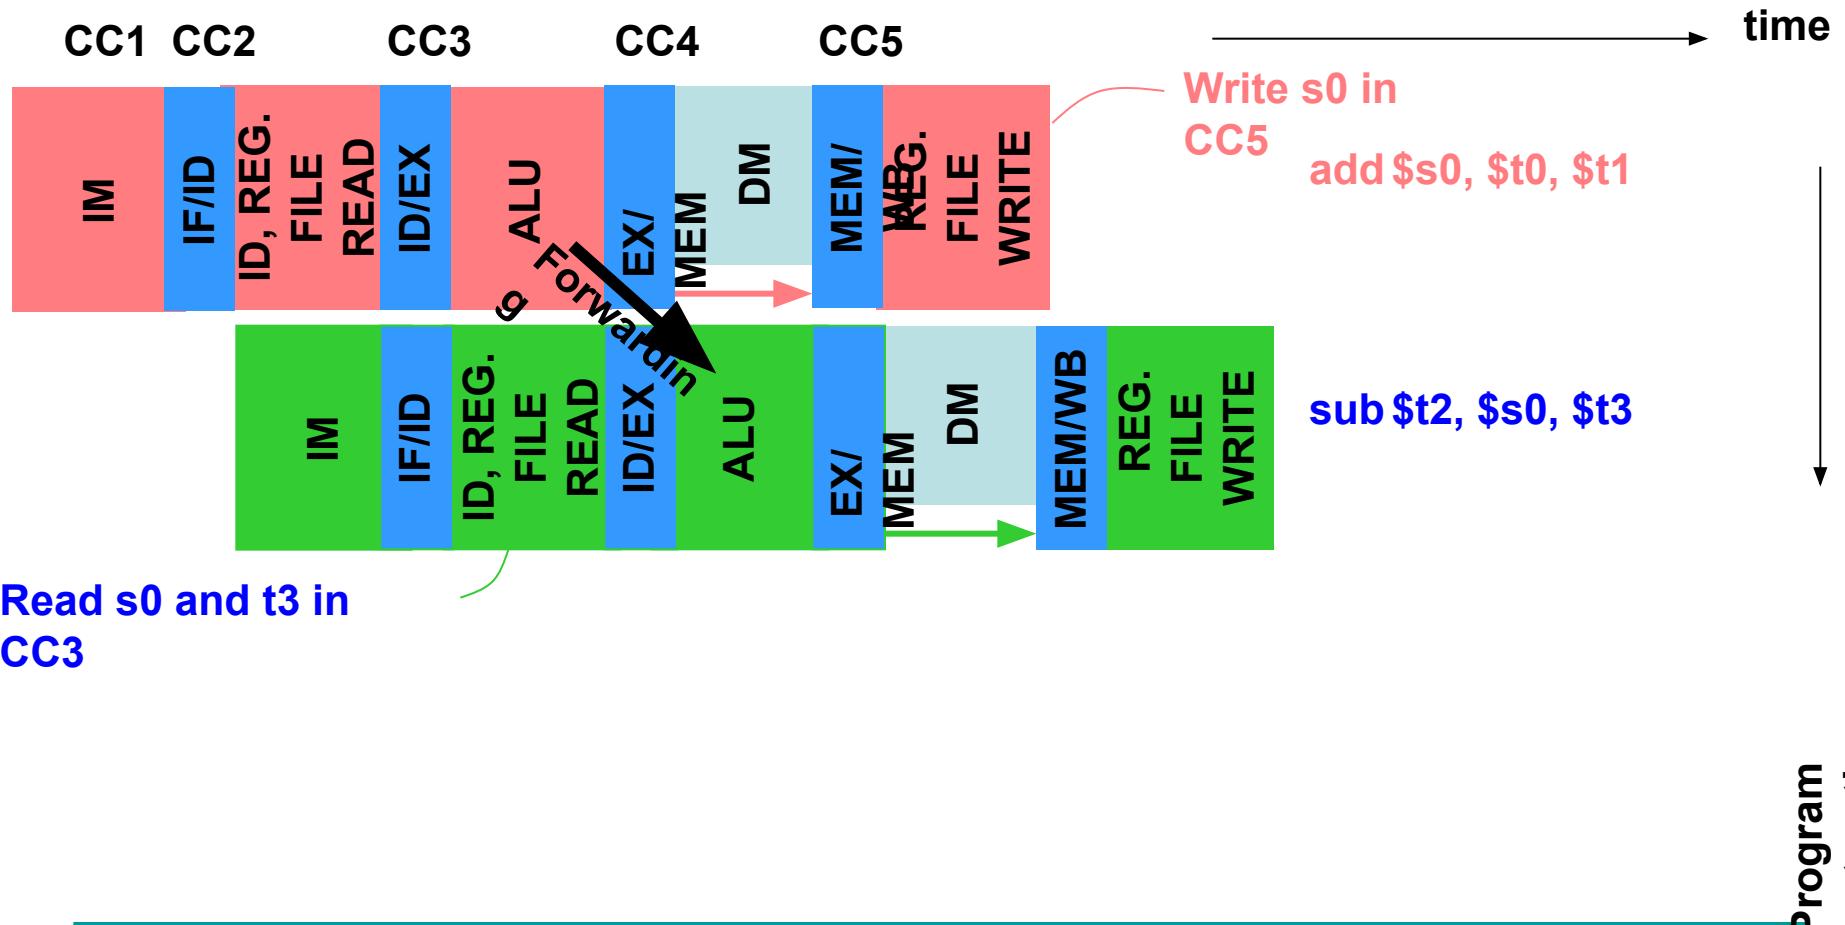


Forwarding or Bypassing

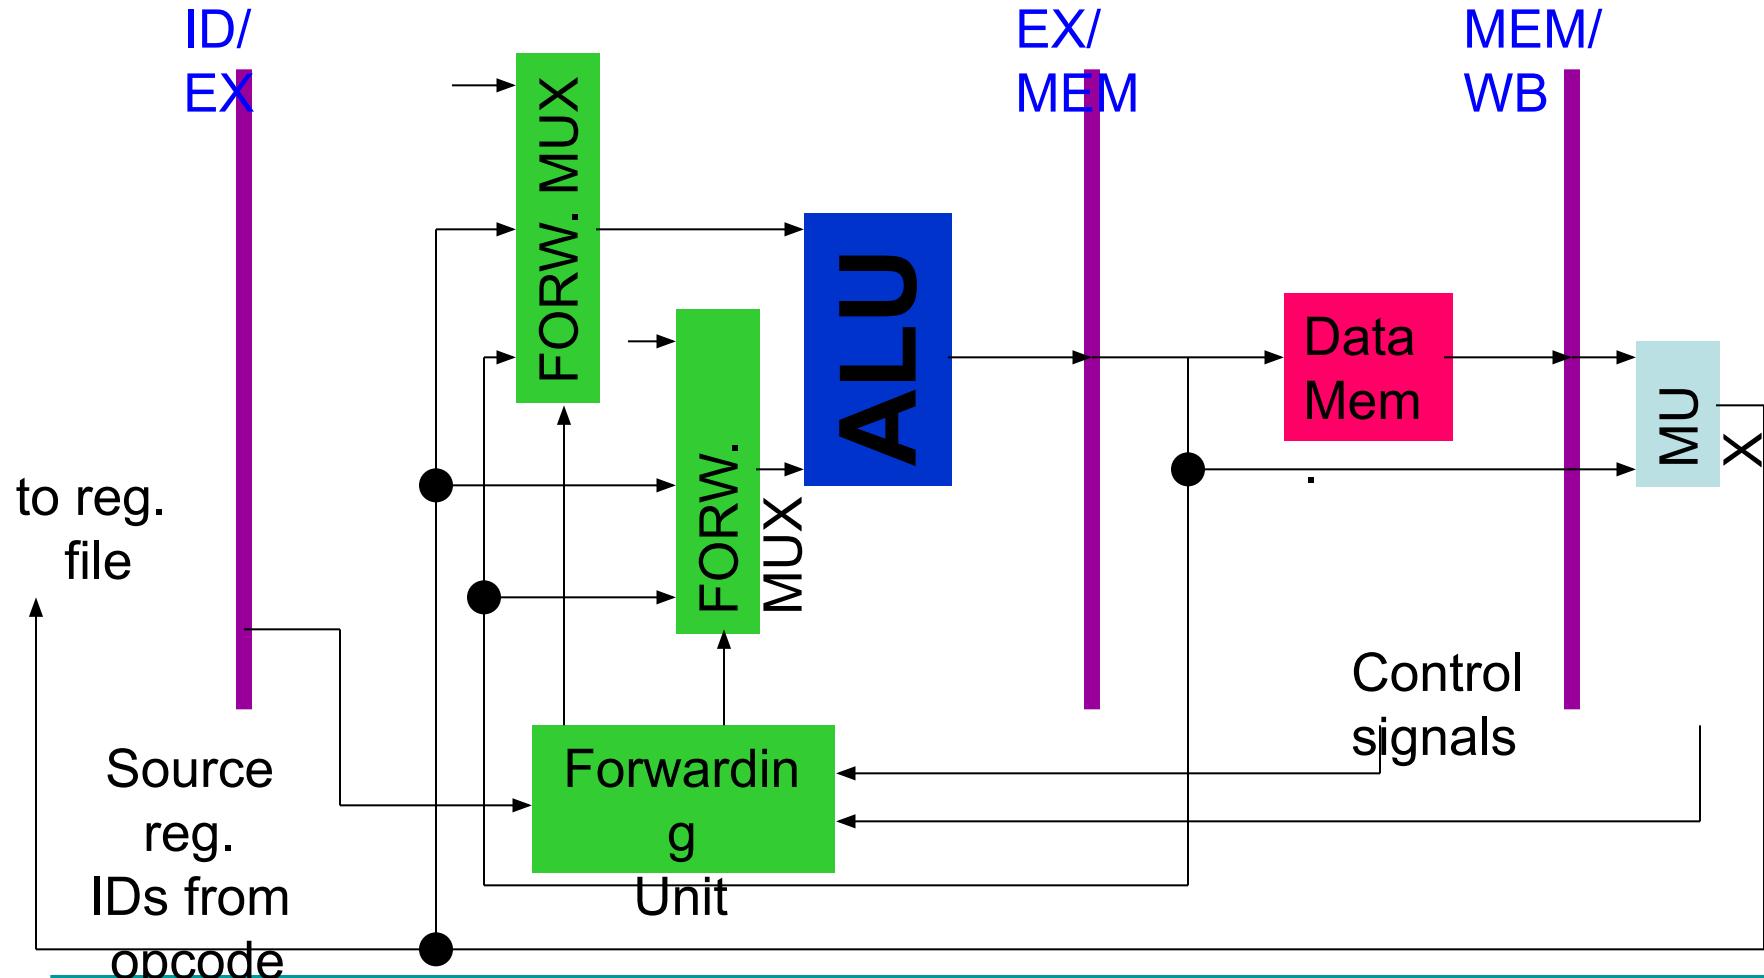
- Output of a resource used by an instruction is forwarded to the input of some resource being used by another instruction.
- Forwarding can eliminate some, but not all, data hazards.



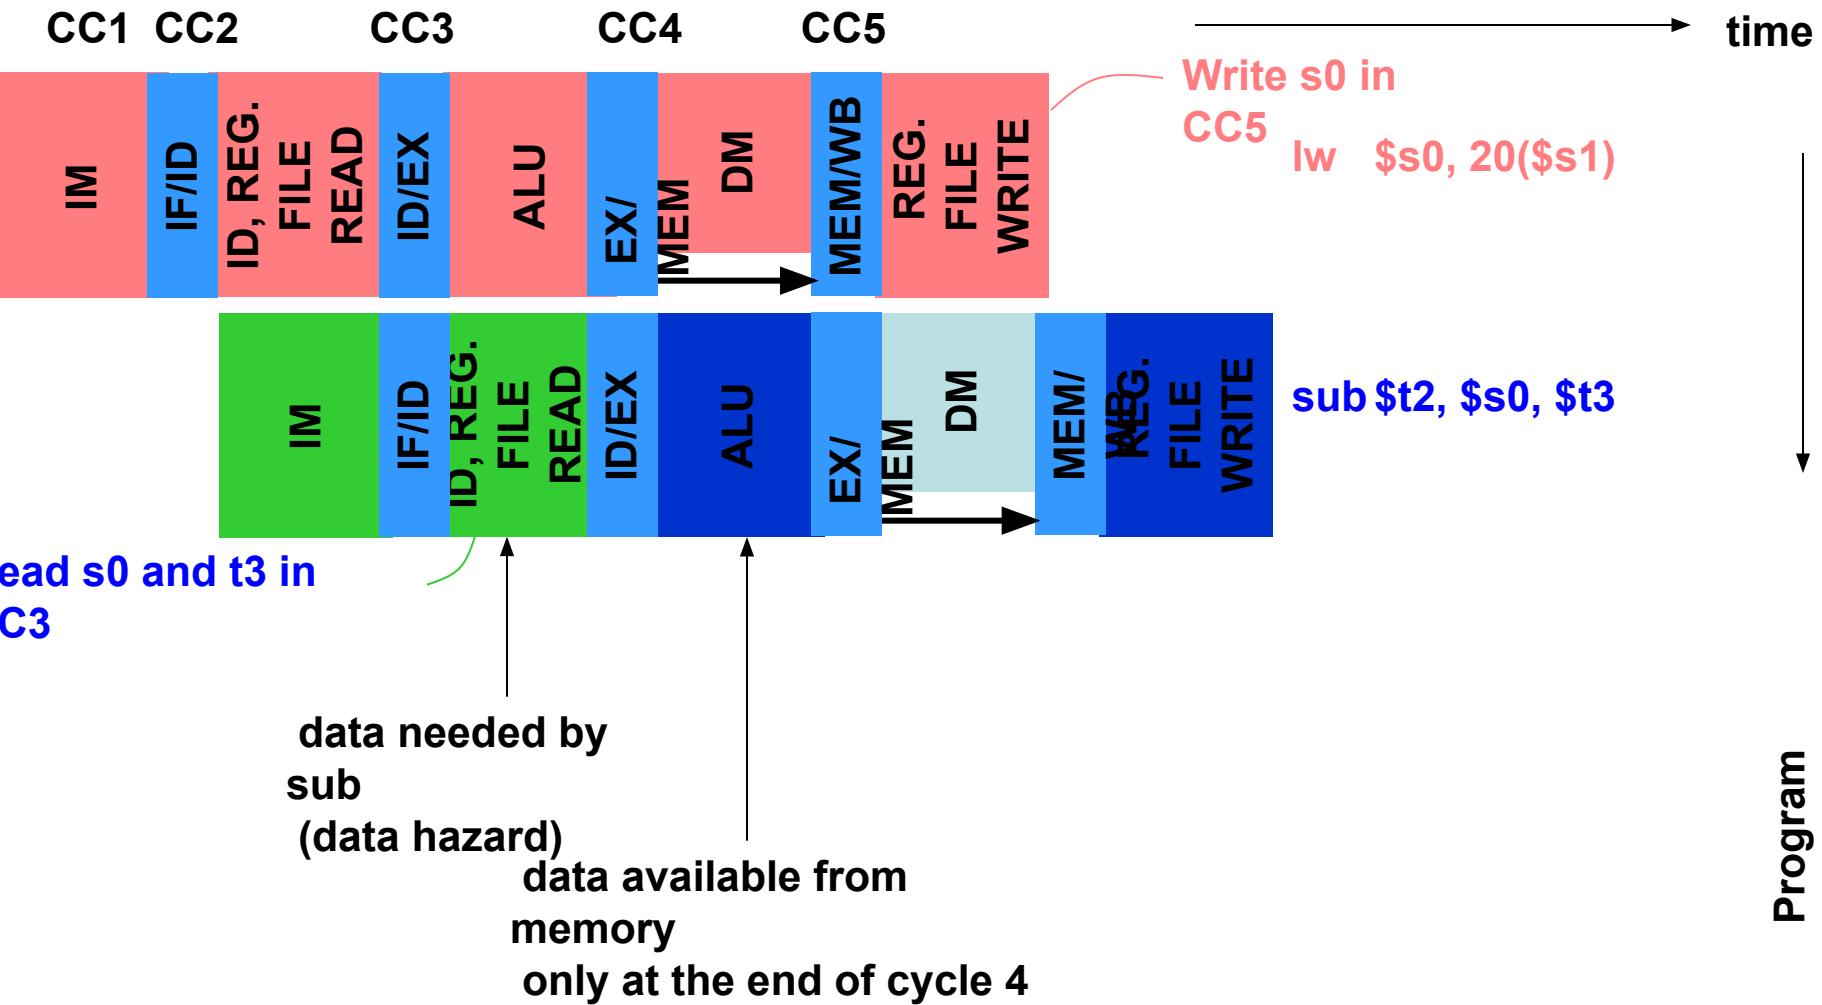
Forwarding for Data Hazard



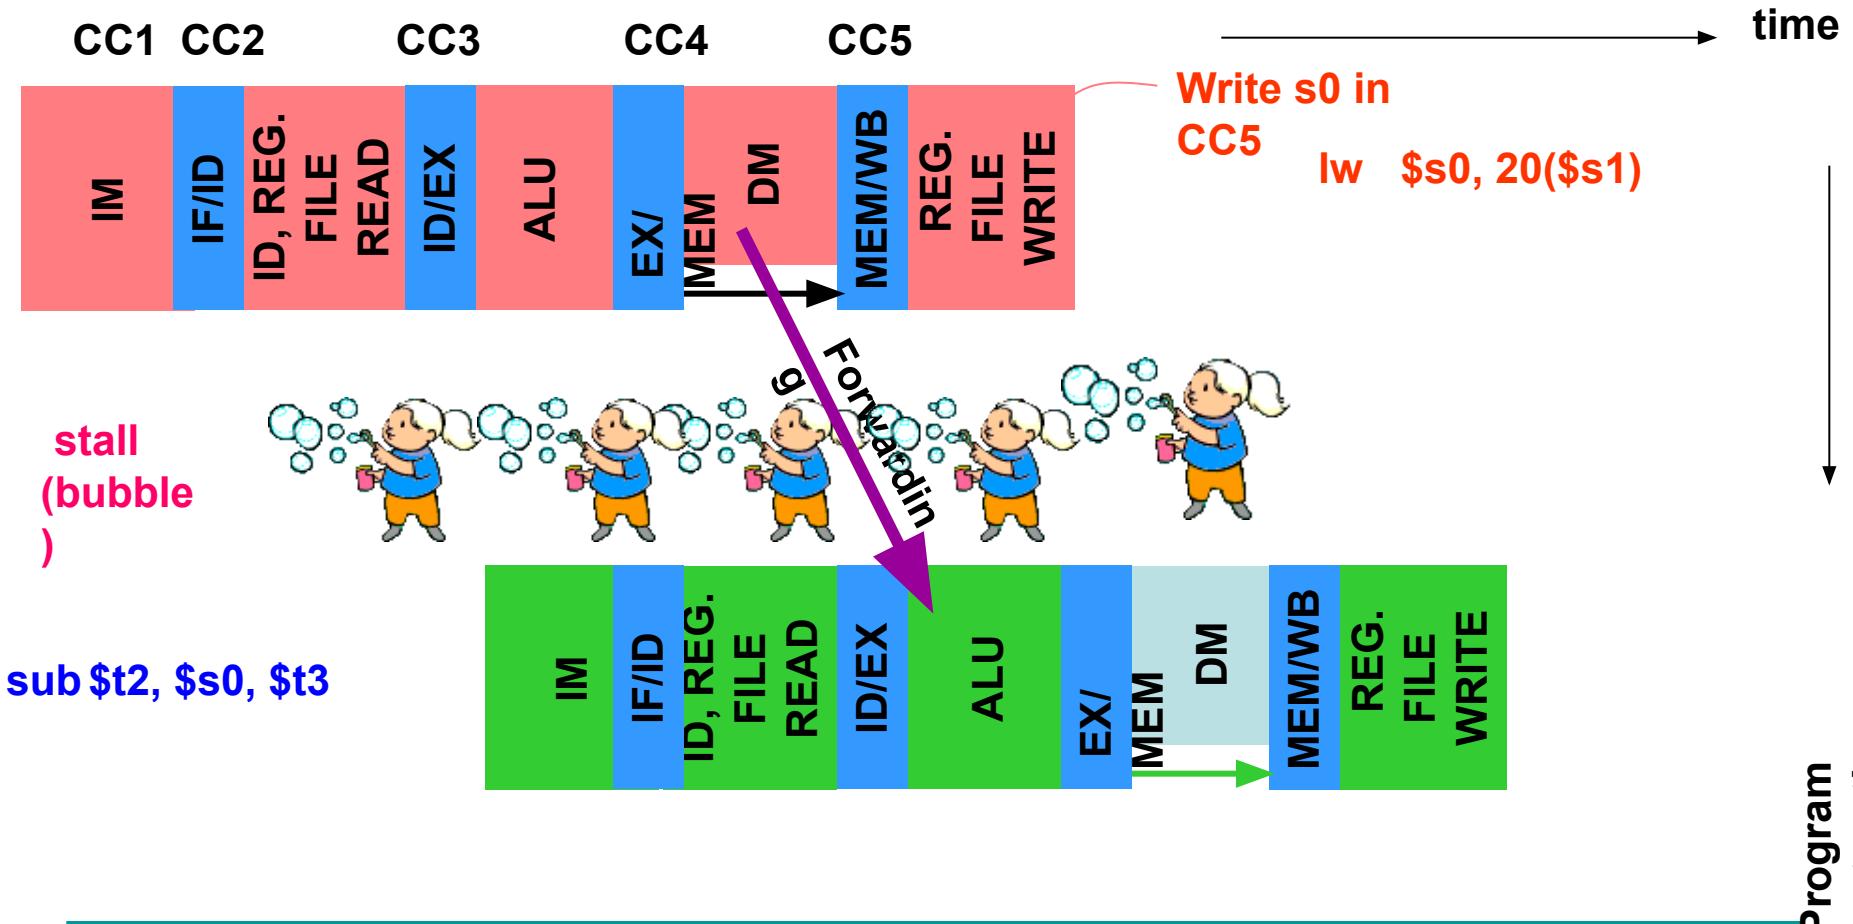
Forwarding Unit Hardware



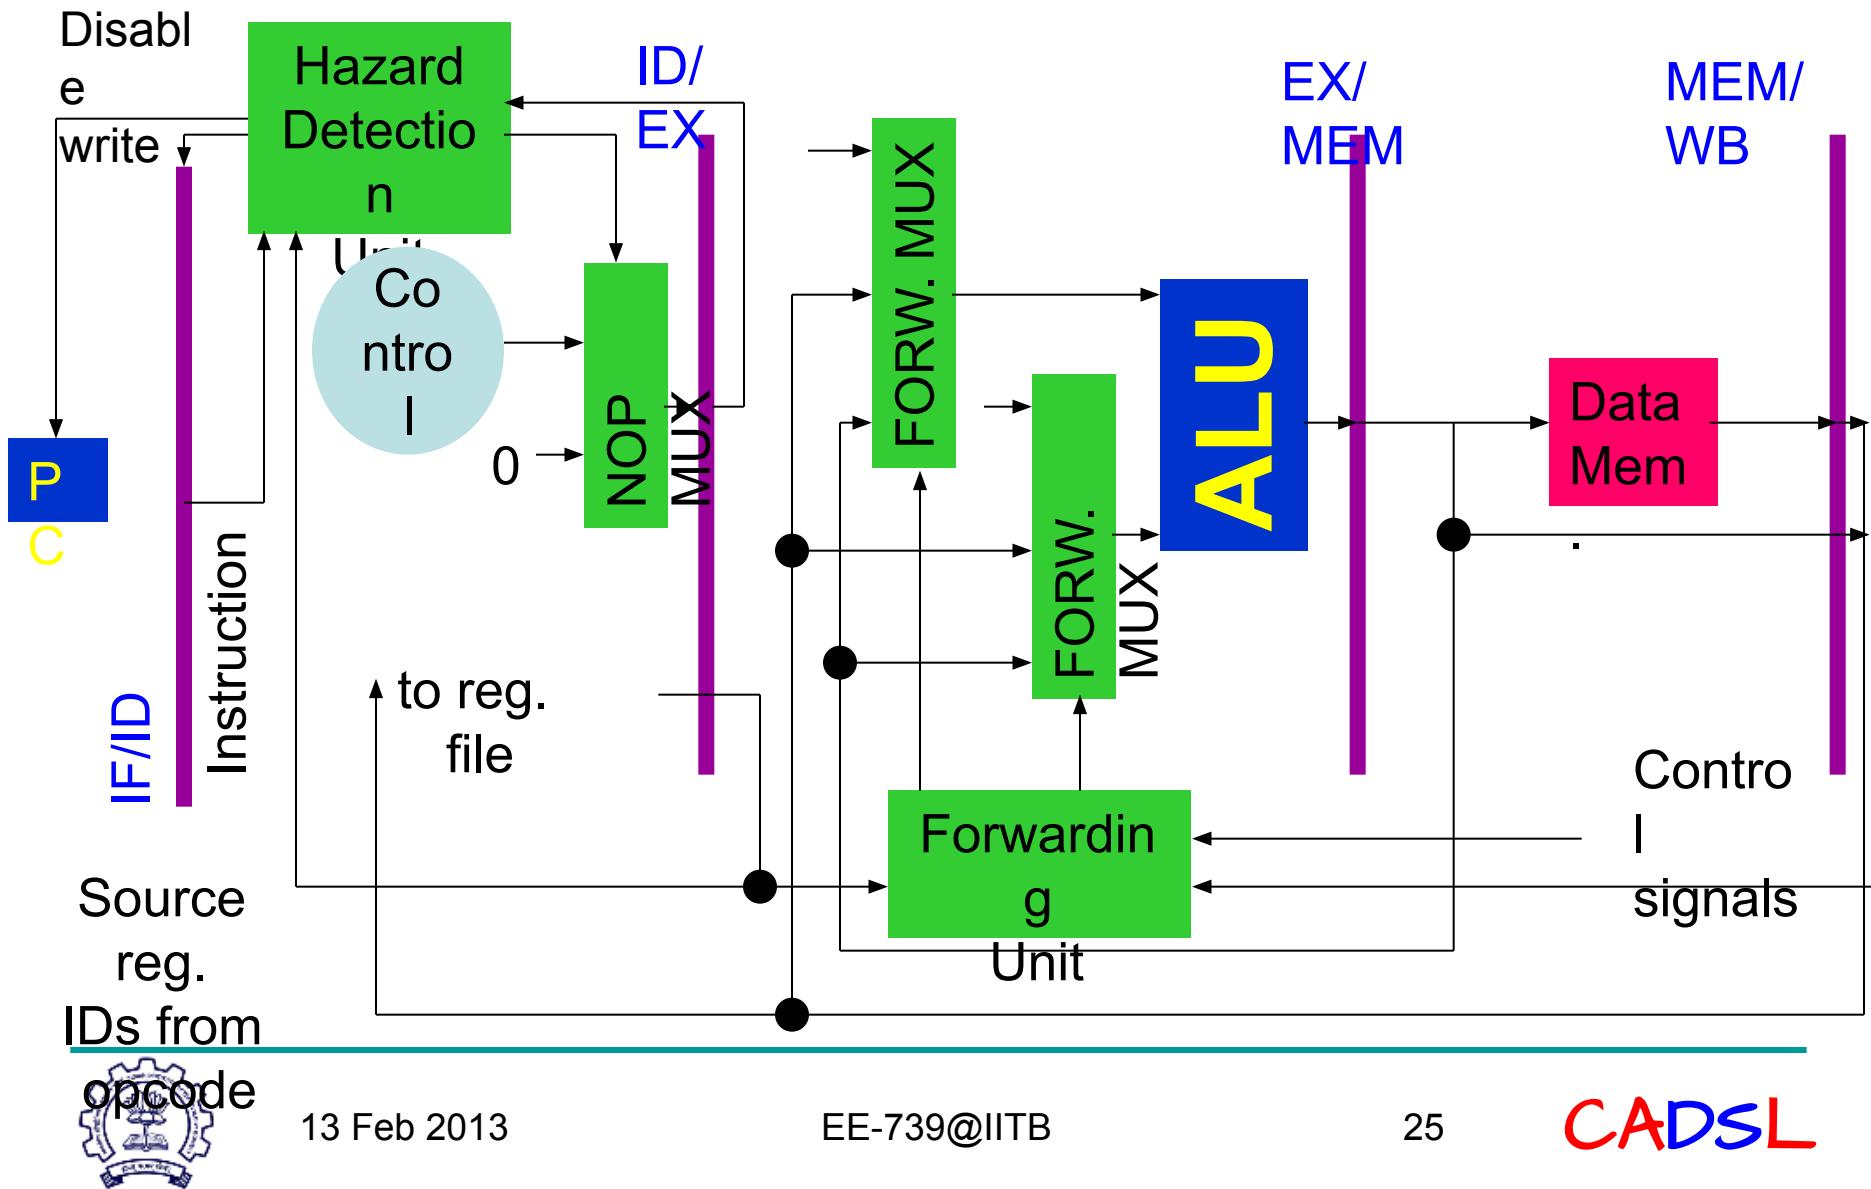
Forwarding Alone May Not Work



Use Bubble and Forwarding



Hazard Detection Unit Hardware



Resolving Hazards

- Hazards are resolved by Hazard detection and forwarding units.
- Compiler's understanding of how these units work can improve performance.



Avoiding Stall by Code Reorder

C code:

A = B + E;

C = B + F;

MIPS code:

lw \$t1,0(\$t0)

lw \$t2,4(\$t0)

add\$t3,\$t1, \$t2

sw \$t3,12(\$t0)

lw \$t4,8(\$t0)

add\$t5,\$t1, \$t4

sw \$t5,16,(\$t0)

\$t1 written

\$t2 written

...

...

...

...

...

...

...

...

\$t1, \$t2 needed

\$t4 written

\$t4 needed



Reordered Code

C code:

A = B + E;

C = B + F;

MIPS code:

lw \$t1,0(\$t0)

lw \$t2,4(\$t0)

lw \$t4,8(\$t0)

~~add\$t3,\$t1, \$t2~~

~~sw \$t3,12(\$t0)~~

~~add\$t5,\$t1, \$t4~~

~~sw \$t5,16,(\$t0)~~

no hazard

no hazard



Control Hazard

- Instruction to be fetched is not known!
- Example: Instruction being executed is branch-type, which will determine the next instruction:

add \$4, \$5, \$6

beq \$1, \$2, 40

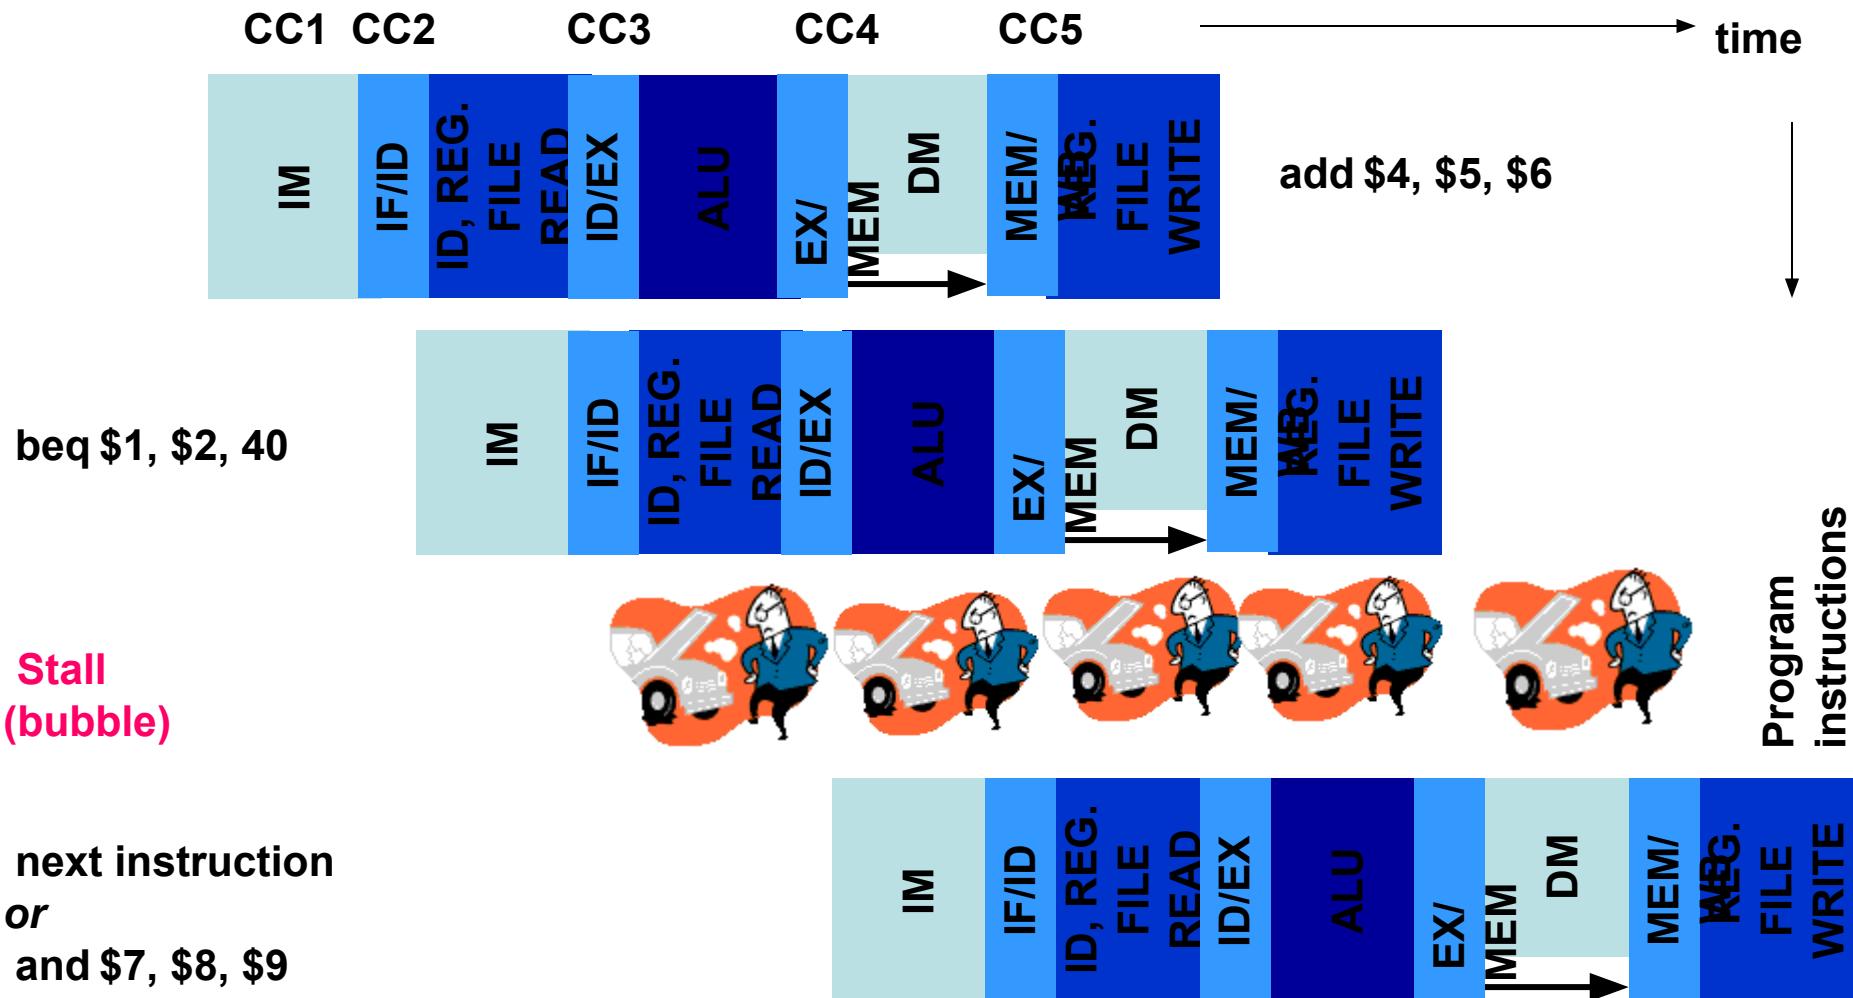
next instruction

...

40and \$7, \$8, \$9



Stall on Branch



Why Only One Stall?

- Extra hardware in ID phase:
 - Additional ALU to compute branch address
 - Comparator to generate zero signal
 - Hazard detection unit writes the branch address in PC

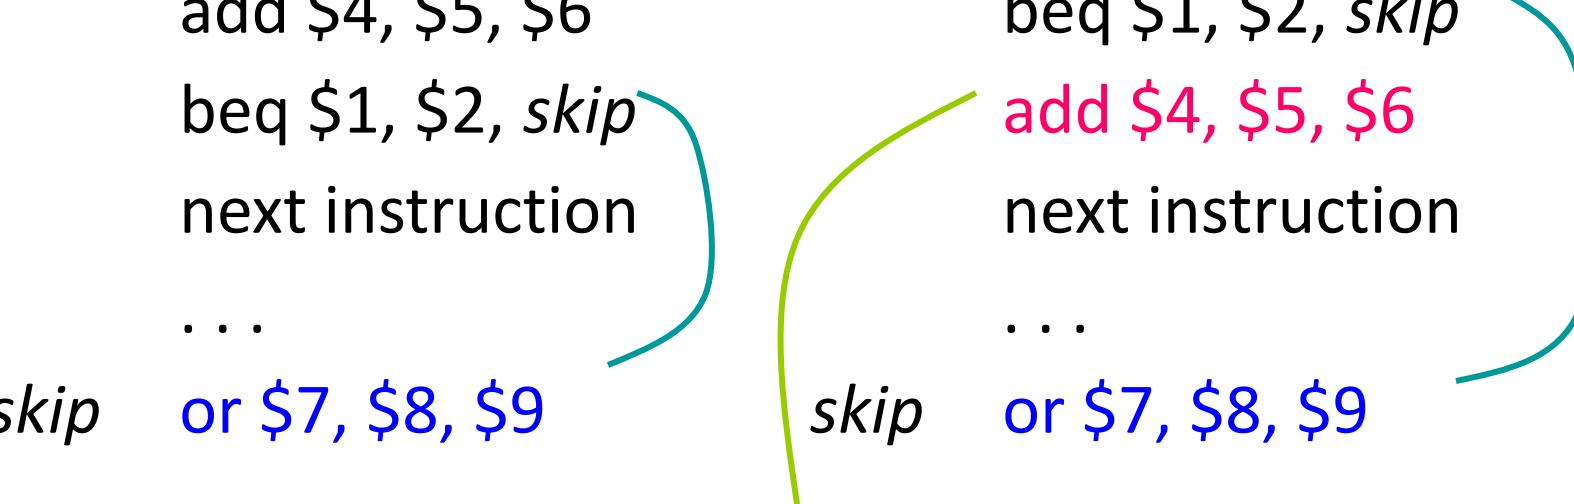


Ways to Handle Branch

- Stall or bubble
- Branch prediction:
 - Heuristics
 - Next instruction
 - Prediction based on statistics (dynamic)
 - Hardware decision (dynamic)
 - Prediction error: pipeline flush
- Delayed branch

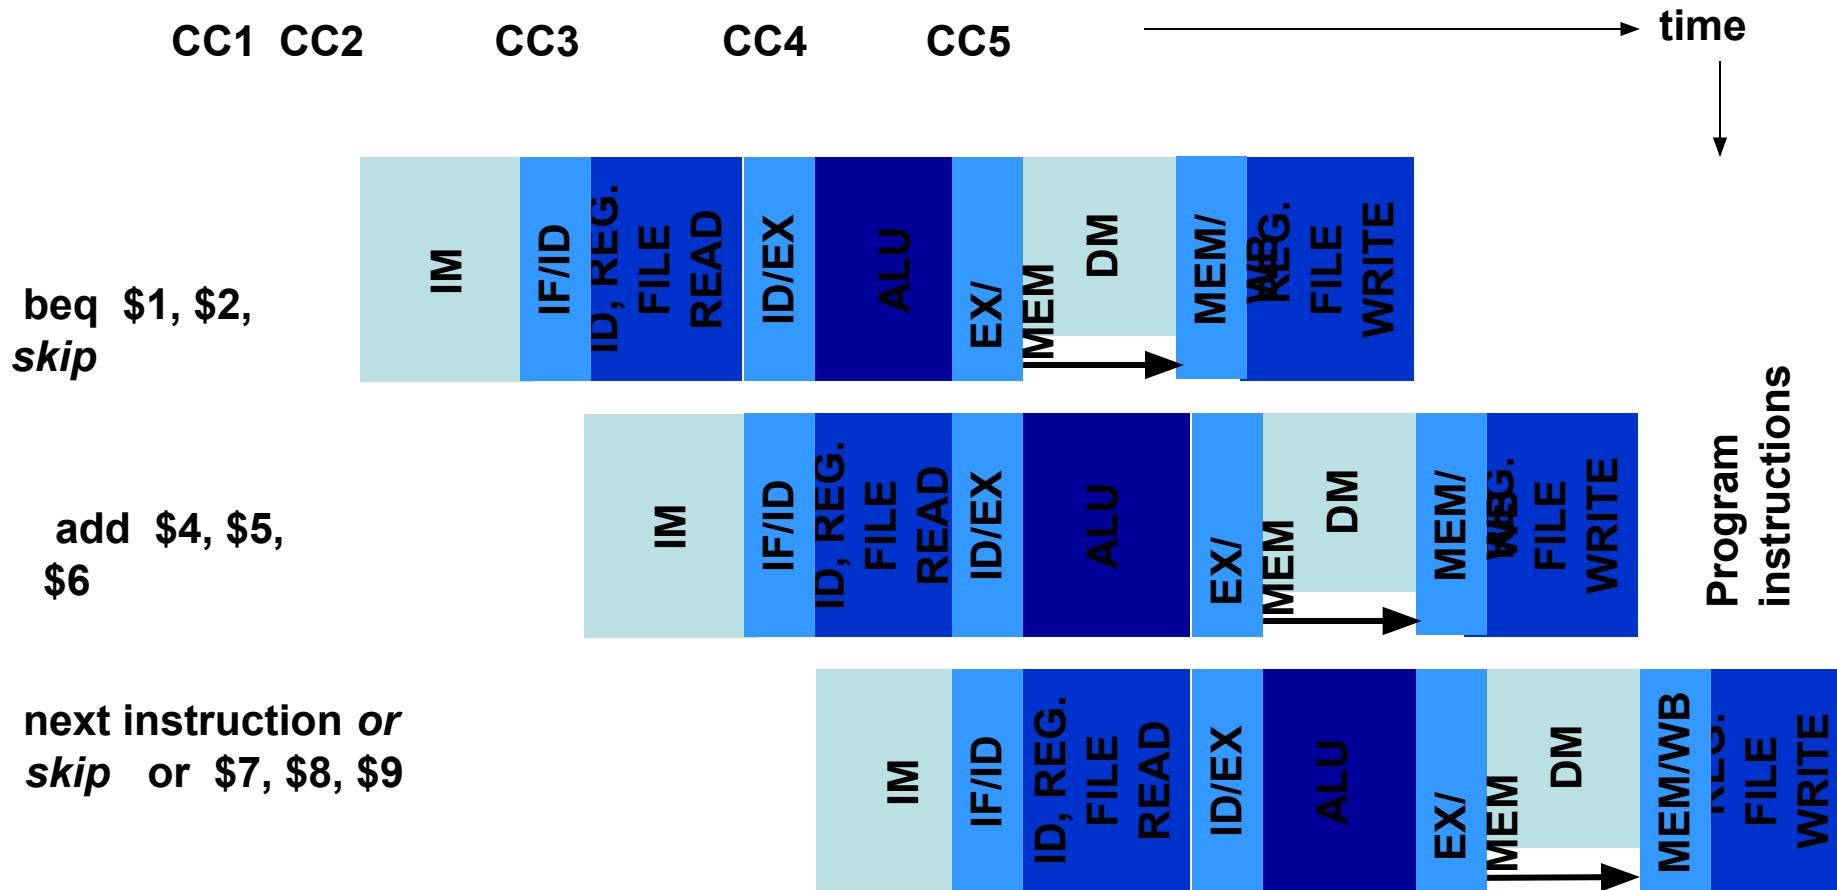


Delayed Branch Example

- Stall on branch
 - add \$4, \$5, \$6
 - beq \$1, \$2, *skip*
 - next instruction
 - ...
 - skip* or \$7, \$8, \$9
 - Delayed branch
 - beq \$1, \$2, *skip*
 - add \$4, \$5, \$6
 - next instruction
 - ...
 - skip* or \$7, \$8, \$9
- Instruction executed irrespective of branch decision*
- 



Delayed Branch

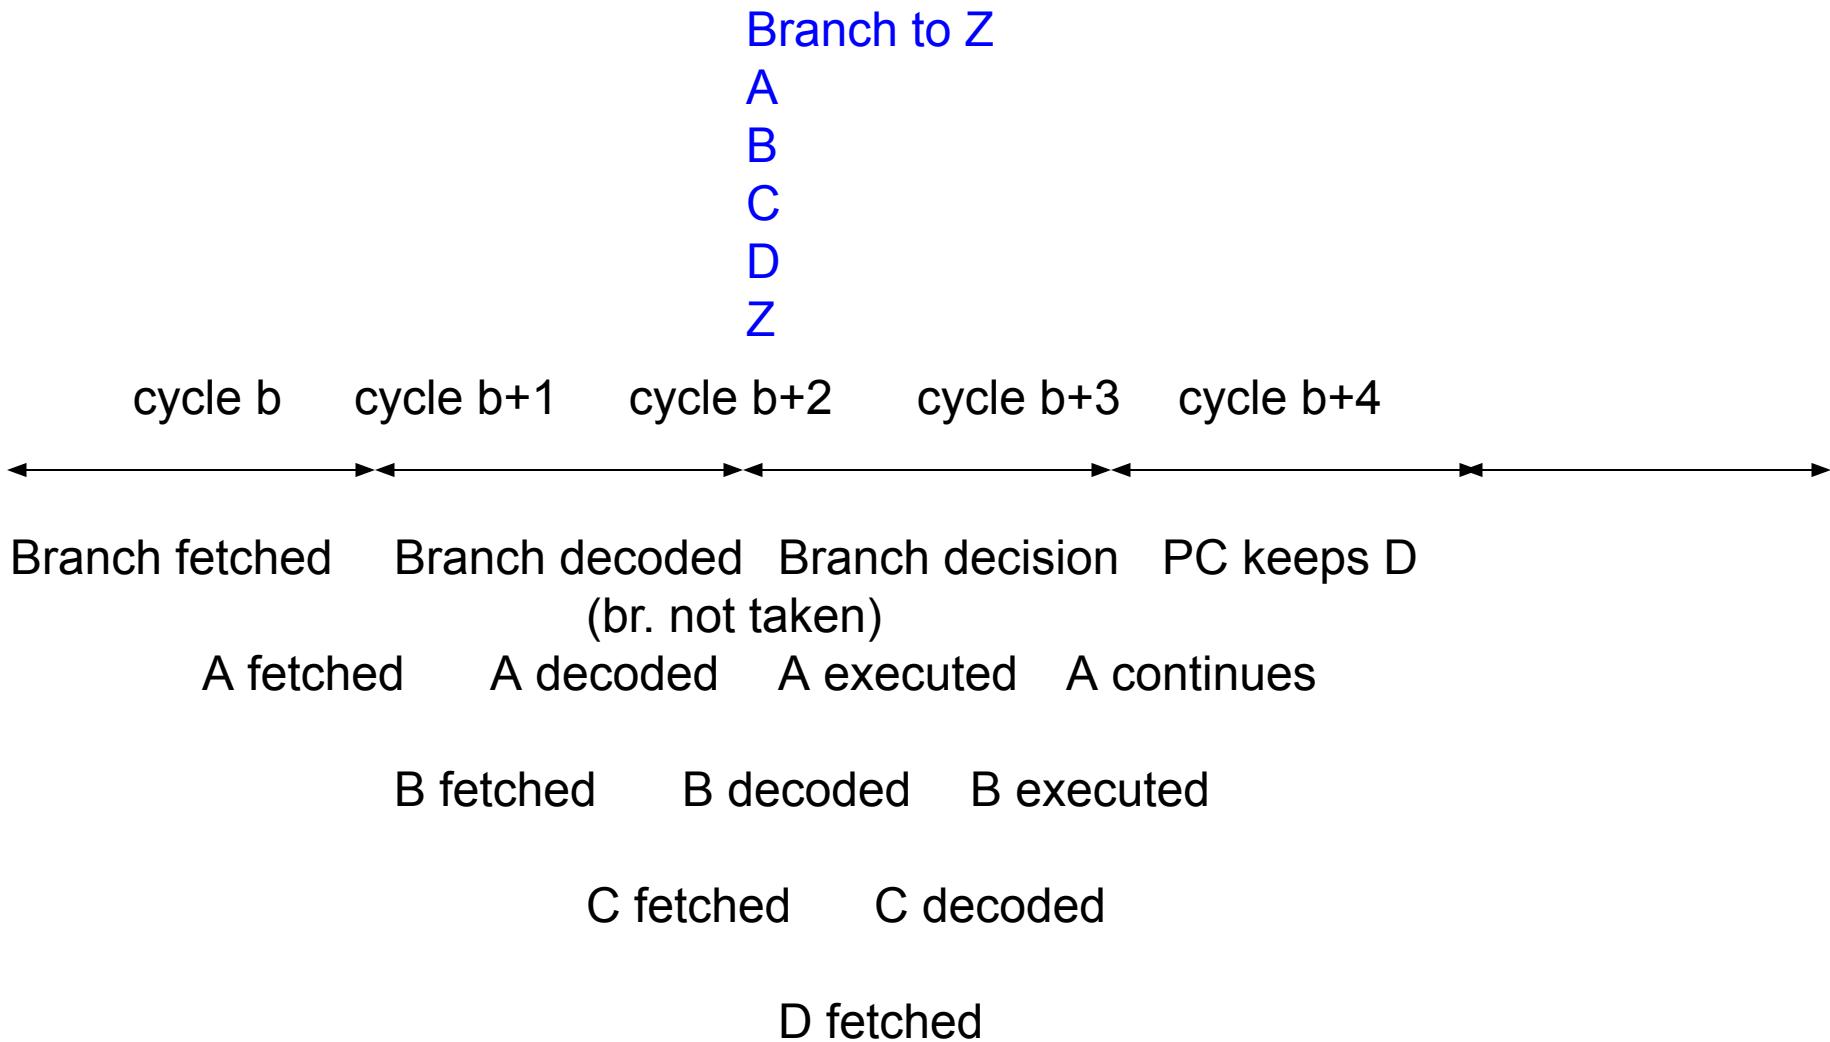


Branch Hazard

- Consider heuristic – branch not taken.
- Continue fetching instructions in sequence following the branch instructions.
- If branch is taken (indicated by *zero* output of ALU):
 - Control generates *branch* signal in ID cycle.
 - *branch* activates *PCSource* signal in the MEM cycle to load PC with new branch address.
 - *Three instructions in the pipeline must be flushed if branch is taken – can this penalty be reduced?*



Branch Not Taken



Branch Taken

Branch to Z

A

B

C

D

Z

cycle b cycle b+1 cycle b+2 cycle b+3 cycle b+4



Branch fetched Branch decoded Branch decision PC gets Z
(br. taken)

A fetched A decoded A executed Nop

B fetched B decoded Nop

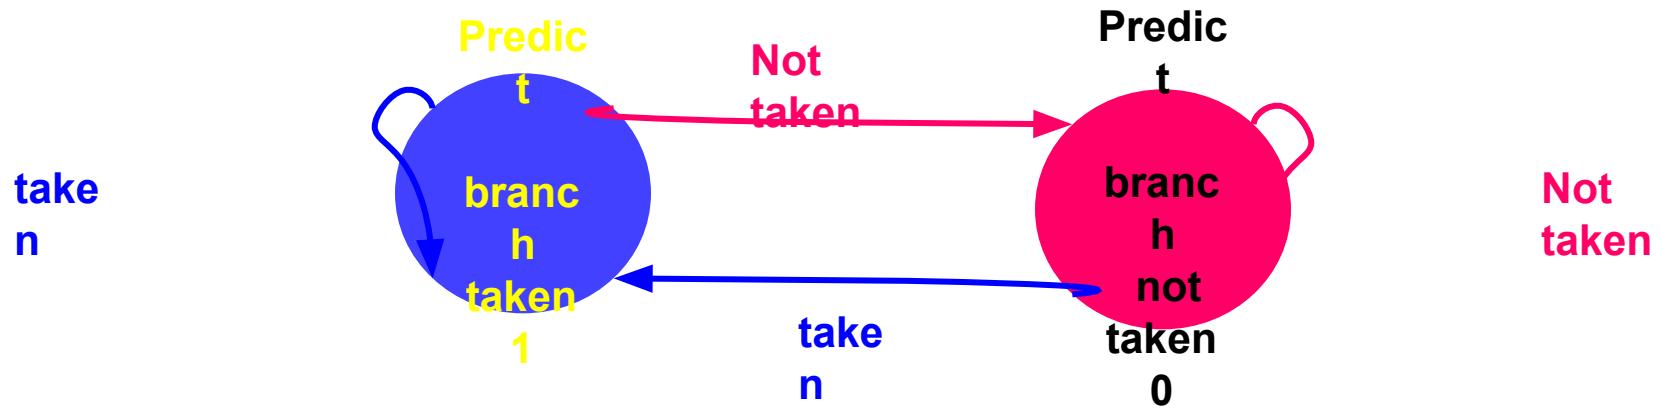
C fetched Nop

*Three instructions are
flushed if branch is
taken*

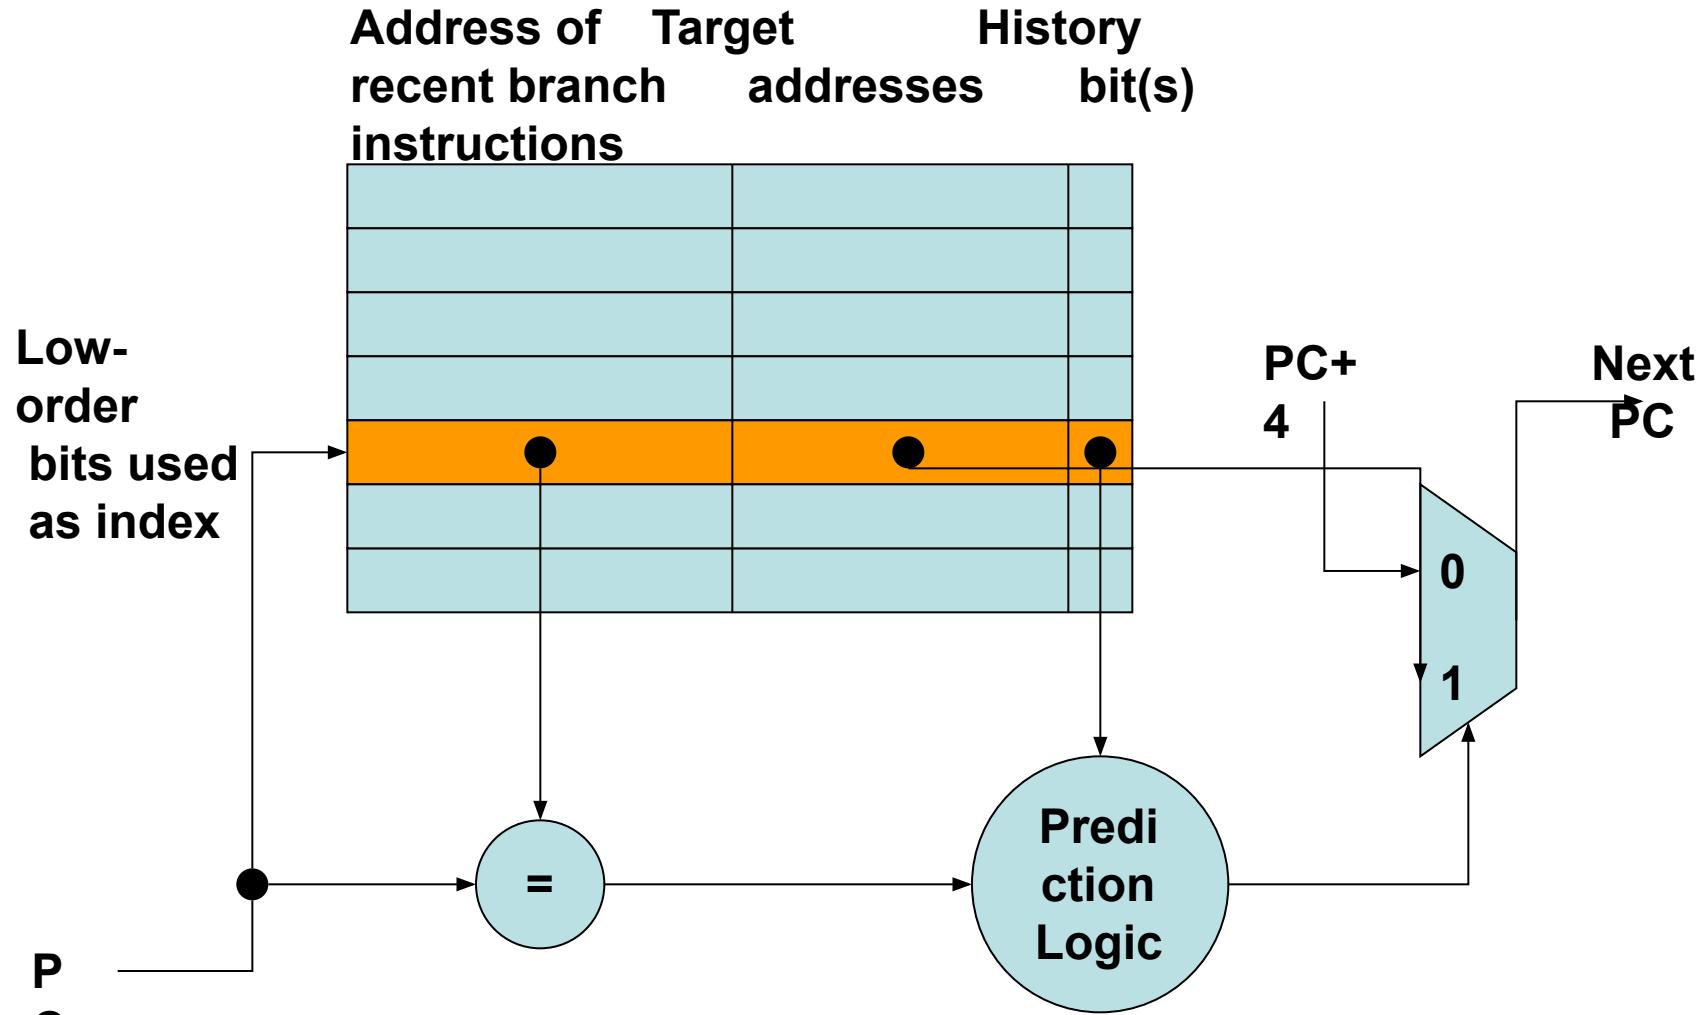


Branch Prediction

- Useful for program loops.
- A one-bit prediction scheme: a one-bit buffer carries a “history bit” that tells what happened on the last branch instruction
 - History bit = 1, branch was taken
 - History bit = 0, branch was not taken

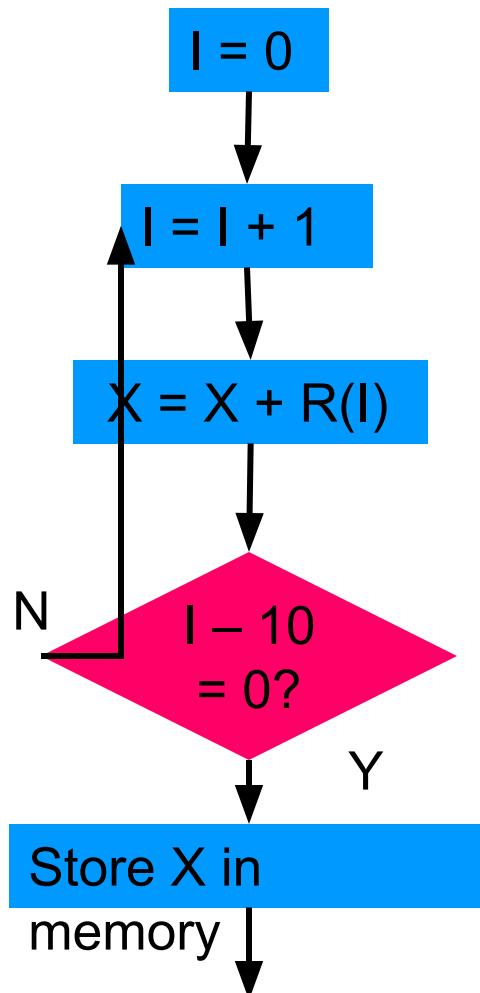


Branch Prediction



Branch Prediction for a Loop

1
2
3
4
5



Execution of Instruction 4

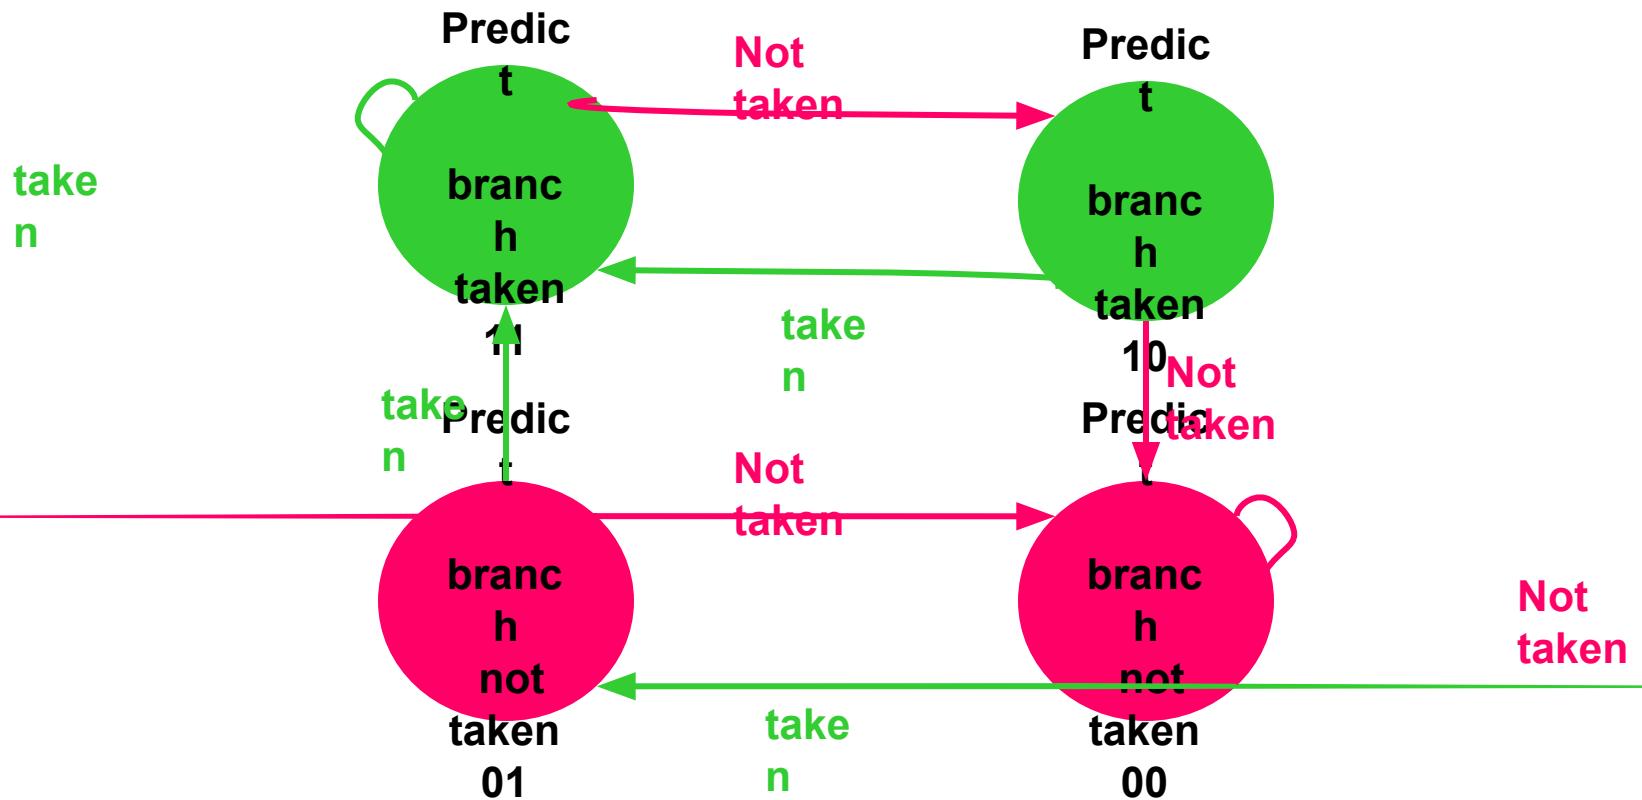
Execu-tion seq.	Old hist. bit	Next instr.			New hist. bit	Prediction
		Pred.	I	Act.		
1	0	5	1	2	1	Bad
2	1	2	2	2	1	Good
3	1	2	3	2	1	Good
4	1	2	4	2	1	Good
5	1	2	5	2	1	Good
6	1	2	6	2	1	Good
7	1	2	7	2	1	Good
8	1	2	8	2	1	Good
9	1	2	9	2	1	Good
10	1	2	10	5	0	Bad

h.bit = 0 branch not taken, h.bit = 1 branch taken.



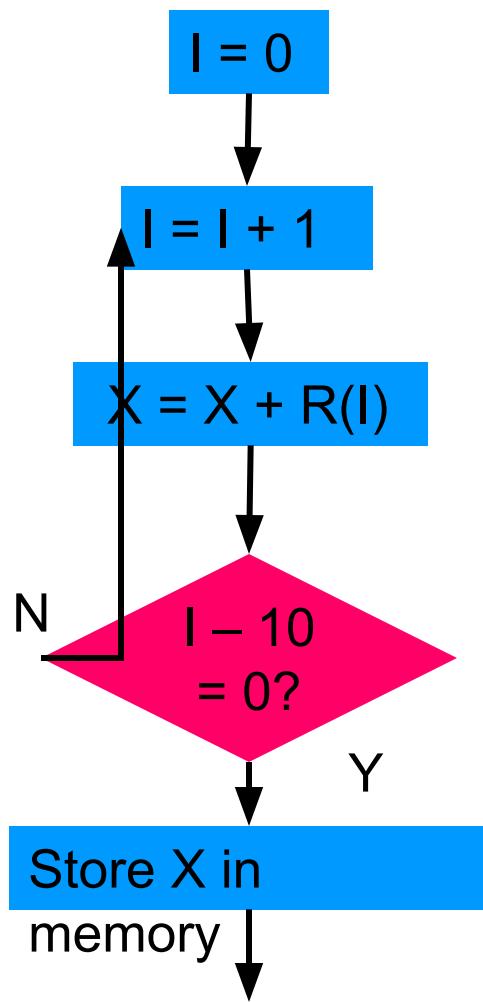
Two-Bit Prediction Buffer

- Can improve correct prediction statistics.



Branch Prediction for a Loop

1
2
3
4
5



Execution of Instruction 4

Execution seq.	Old Pred. Buf	Next instr.			New pred. Buf	Prediction
		Pred.	I	Act.		
1	10	2	1	2	11	Good
2	11	2	2	2	11	Good
3	11	2	3	2	11	Good
4	11	2	4	2	11	Good
5	11	2	5	2	11	Good
6	11	2	6	2	11	Good
7	11	2	7	2	11	Good
8	11	2	8	2	11	Good
9	11	2	9	2	11	Good
10	11	2	10	5	10	Bad



Summary: Hazards

- Structural hazards
 - Cause: resource conflict
 - Remedies: (i) hardware resources, (ii) stall (bubble)
- Data hazards
 - Cause: data unavailability
 - Remedies: (i) forwarding, (ii) stall (bubble), (iii) code reordering
- Control hazards
 - Cause: out-of-sequence execution (branch or jump)
 - Remedies: (i) stall (bubble), (ii) branch prediction/pipeline flush, (iii) delayed branch/pipeline flush



Limits of Pipelining

- IBM RISC Experience
 - Control and data dependences add 15%
 - Best case CPI of 1.15, IPC of 0.87
 - Deeper pipelines (higher frequency) magnify dependence penalties
- This analysis assumes 100% cache hit rates
 - Hit rates approach 100% for some programs
 - Many important programs have much worse hit rates



Processor Performance

$$\text{Processor Performance} = \frac{\text{Program Time}}{\text{Instructions}}$$
$$= \frac{\text{Program code size}}{\text{Cycles per Instruction}} \times \frac{\text{Time per cycle}}$$

Program
Instructions
Cycles per Instruction
Time per cycle

- In the 1980's (decade of pipelining):
 - CPI: 5.0 => 1.15
- In the 1990's (decade of superscalar):
 - CPI: 1.15 => 0.5 (best case)
- In the 2000's (decade of multicore):
 - Marginal CPI improvement

