

# Multi Level Feedback Queue (MLFQ) Scheduling

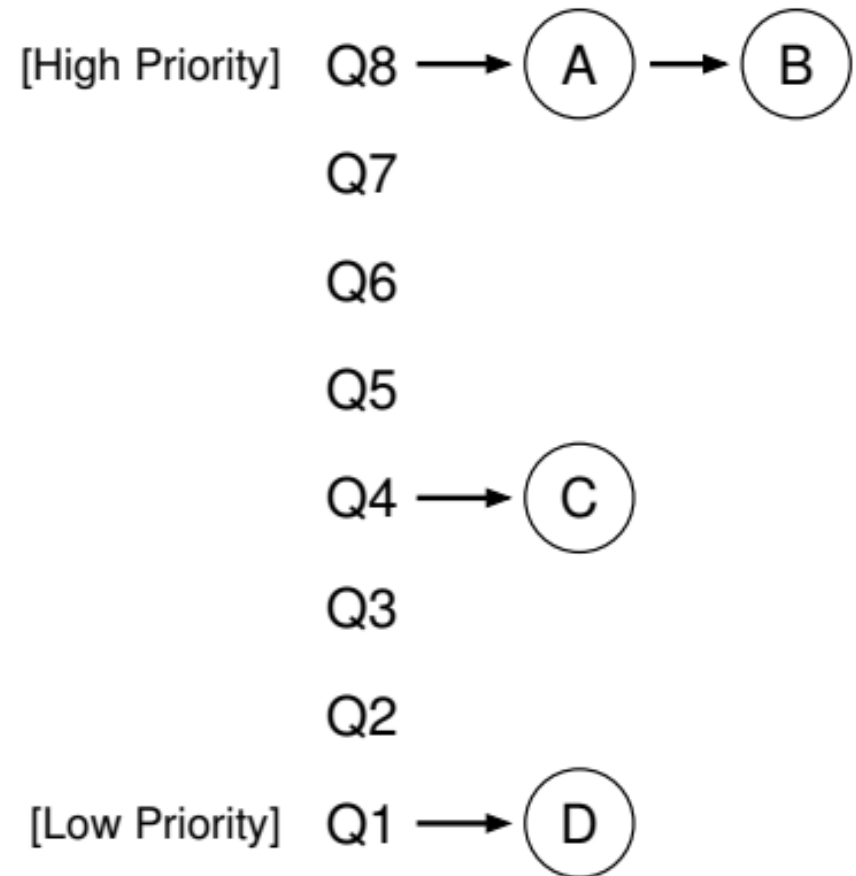
- optimize turnaround time
- minimize response time - system feel responsive to interactive users
- number of distinct queues each assigned a different priority level
- a job that is ready to run is on a single queue

first two basic rules for MLFQ:

Rule 1: If  $\text{Priority}(A) > \text{Priority}(B)$ , A runs (B doesn't)

Rule 2: If  $\text{Priority}(A) = \text{Priority}(B)$ , A & B run in RR.


- MLFQ varies the priority of a job based on its observed behavior
  - If a job repeatedly relinquishes the CPU while waiting for input from the keyboard, **MLFQ will keep its priority high**, as this is how an interactive process might behave
  - If, instead, a job uses the CPU intensively for long periods of time, **MLFQ will reduce its priority**




# Changing Priority


- **Rule 3:** When a job enters the system, it is placed at the highest priority (the topmost queue)
- **Rule 4a:** If a job uses up an entire time slice while running, its priority is *reduced* (i.e., it moves down one queue).
- **Rule 4b:** If a job gives up the CPU before the time slice is up, it stays at the *same* priority level.

Long Process A (CPU intensive)



Short running interactive Process B






- Mixture of I/O-intensive and CPU-intensive Workload
- Interactive Process B (gray) that uses CPU for only 1 ms

# Priority Boost


- Move all the jobs to the topmost queue after time period.
- Advantages:
  - processes are guaranteed not to starve



- **Rule 5:** After some time period  $S$ , move all the jobs in the system to the topmost queue.

What is the value for  $S$ ?


# Gaming the scheduler



- generally refers to the idea of doing something sneaky to trick the scheduler into giving you more than your fair share of the resource
- before the time slice is over, issue an I/O operation (to some file you don't care about) and thus relinquish the CPU; doing so allows you to remain in the same queue, and thus gain a higher percentage of CPU time
- thereby, a job could nearly monopolize the CPU

- **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down one queue)

# Lower Priority, Longer Quanta





# Summary - MLFQ

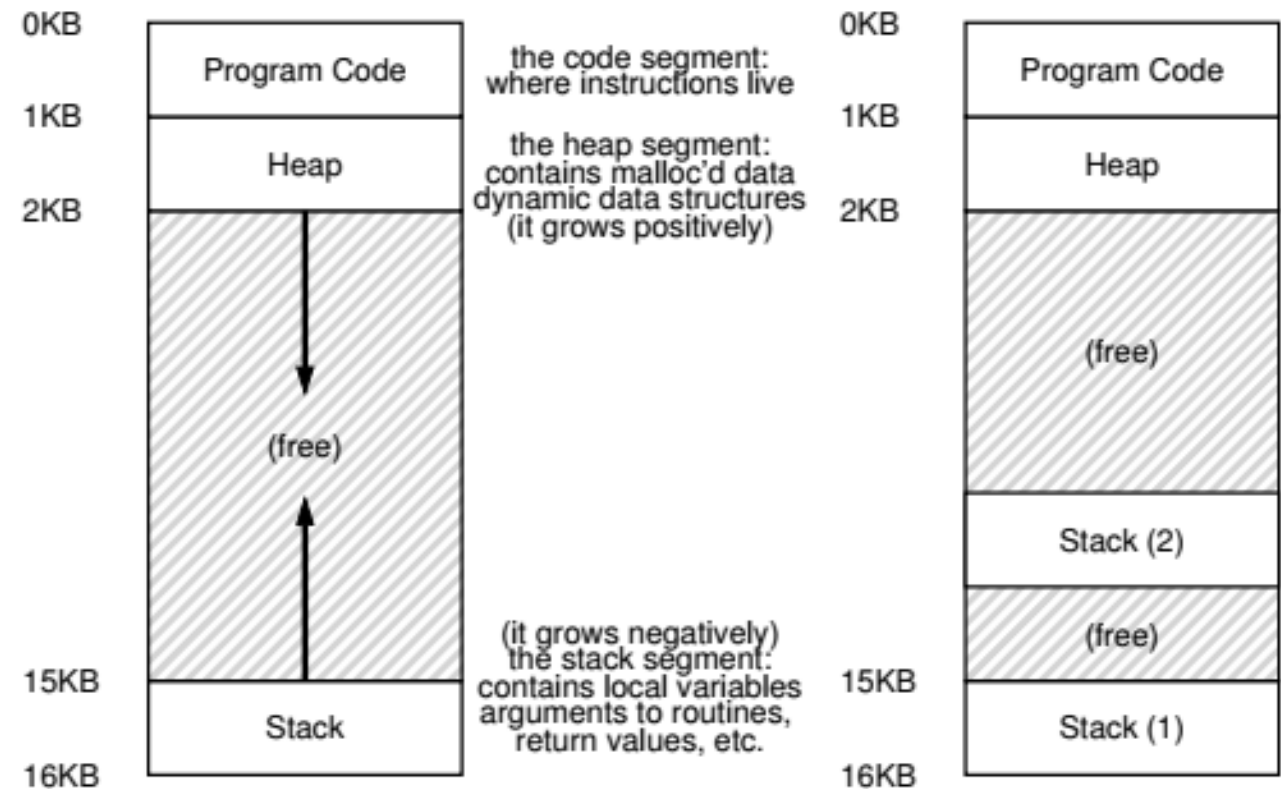
- **Rule 1:** If  $\text{Priority}(A) > \text{Priority}(B)$ , A runs (B doesn't).
- **Rule 2:** If  $\text{Priority}(A) = \text{Priority}(B)$ , A & B run in round-robin fashion using the time slice (quantum length) of the given queue.
- **Rule 3:** When a job enters the system, it is placed at the highest priority (the topmost queue).
- **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down one queue).
- **Rule 5:** After some time period  $S$ , move all the jobs in the system to the topmost queue.

# Threads

- Concept of Process
  - **Resource ownership:** A process includes a virtual address space to hold the process image, i.e., the collection of program, data, stack, and attributes defined in the PCB. Resources are main memory, Disk I/O, I/O devices, and files
  - **Scheduling/execution:** The execution of a process follows an execution path (trace) through one or more programs. A process has an execution state (Running, Ready, etc.) and a dispatching priority
- Threading
  - Ability of an OS to support multiple, concurrent paths of execution within a single process.
  - Lightweight process
  - Achieves parallelism

# Threads ... Contd.

- Single point of execution within a program
- Share same address space and can access same data
- Context switching: between threads; Adv.: address space remains the same, page table too
- Supports parallelism



Single-Threaded and Multi-Threaded Address Spaces

```

1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4  #include "common.h"
5  #include "common_threads.h"
6
7  void *mythread(void *arg) {
8      printf("%s\n", (char *) arg);
9      return NULL;
10 }
11
12 int
13 main(int argc, char *argv[]) {
14     pthread_t p1, p2;
15     int rc;
16     printf("main: begin\n");
17     Pthread_create(&p1, NULL, mythread, "A");
18     Pthread_create(&p2, NULL, mythread, "B");
19     // join waits for the threads to finish
20     Pthread_join(p1, NULL);
21     Pthread_join(p2, NULL);
22     printf("main: end\n");
23     return 0;
24 }

```

- Main program creates two threads
- `pthread_create()` : Creates thread
- `Pthread_join()`: waits for a particular thread to complete

## Thread trace (1)

main	Thread 1	Thread2
starts running prints "main: begin" creates Thread 1 creates Thread 2 waits for T1	runs prints "A" returns	
waits for T2		runs prints "B" returns
prints "main: end"		

## Thread trace (2)


main	Thread 1	Thread2
starts running prints "main: begin" creates Thread 1	runs prints "A" returns	
creates Thread 2		runs prints "B" returns
waits for T1 <i>returns immediately; T1 is done</i> waits for T2 <i>returns immediately; T2 is done</i> prints "main: end"		

## Thread trace (3)

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
		runs
		prints "B"
		returns
waits for T1		
	runs	
	prints "A"	
	returns	
waits for T2		
<i>returns immediately; T2 is done</i>		
prints "main: end"		

# Threads

- Four threads created
- Each thread is independent
- Management of threads is simpler than processes
- Shared instructions, global, and heap regions
- Each thread has its own stack




- **Process:**

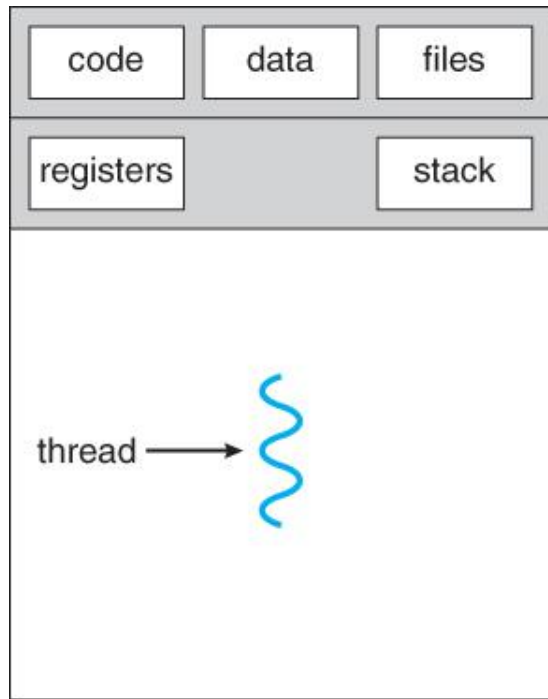
- A virtual address space that holds the process image
- Protected access to processors, other processes, files, and I/O resources

- **Threads**

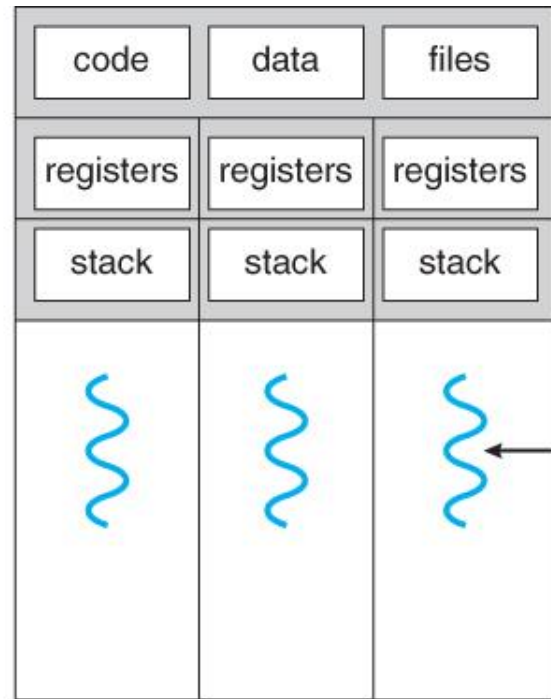
- A thread execution state (Running, Ready, etc.)
- A saved thread context when not running; one way to view a thread is as a independent program counter operating within a process
- An execution stack
- Some per-thread static storage for local variables
- Access to the memory and resources of its process, shared with all other threads in that process







single-threaded process



multithreaded process

Per process	Per thread
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Signals and signal handlers	
Accounting info	

# POSIX threads – IEEE 1003.1c

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NTHREADS 10
void *print_hello_world(void *tid)
{
    /* This function prints the thread's
    identifier and then exits. */
    printf("Hello World. Greetings from
    thread %d\n", tid);
    pthread_exit(NULL);
}
```


```
int main (int argc, char *argv[]) {
    /* The main program creates 10 threads and then exits. */
    pthread_t threads[NTHREADS];
    int status, i;
    for(i=0; i < NTHREADS; i++) {
        printf("Main here. Creating thread %d\n", i);
        status = pthread_create(&threads[i], NULL, print_hello_world,
        (void *)i);
        if (status != 0) {
            printf("pthread returned error code %d\n", status);
            exit(-1);
        }
    }
    exit(NULL);
}
```

# Sum of first 1,00,00,000 numbers

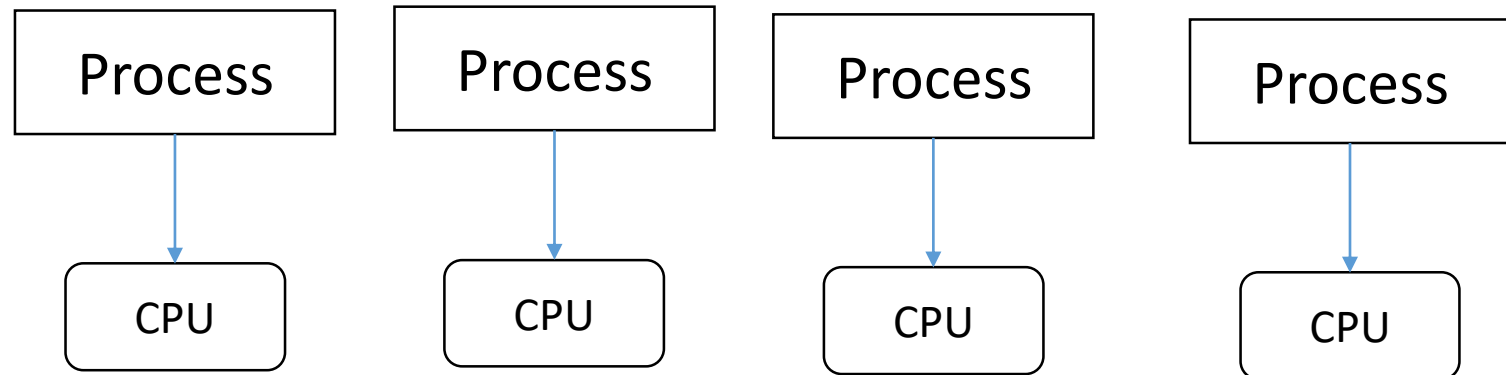
```
#include<stdio.h>
long add() {
    int i=0;
    long sum=0;

    while(i < 10000000) {
        sum = sum+=i;
        i++;
    }
    return sum;
}
```

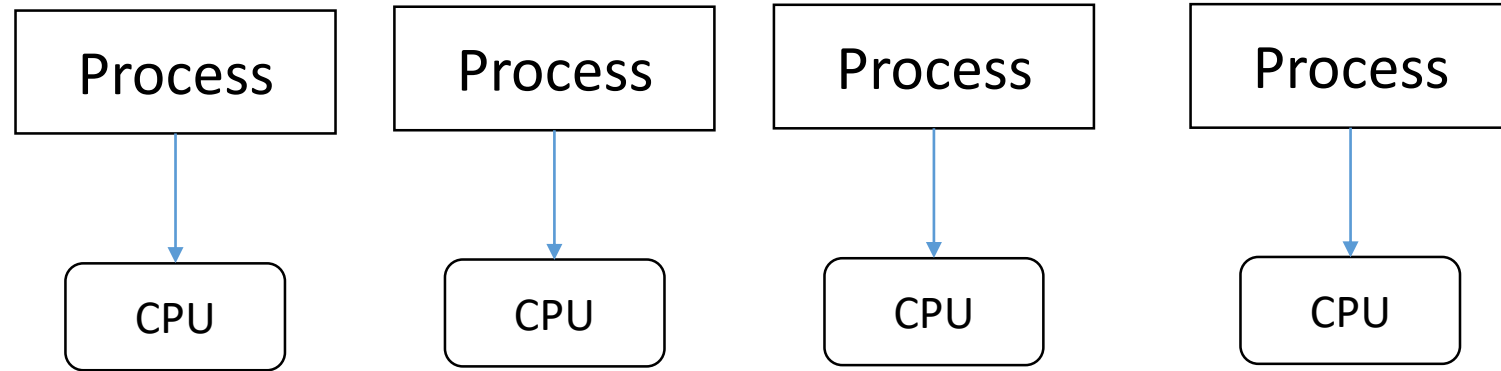
```
int main() {
    long sum;
    sum = add();
    printf("%l", sum);
}
```



Is it possible to speed up the operation? How?  
Assume multiple processors/CPU

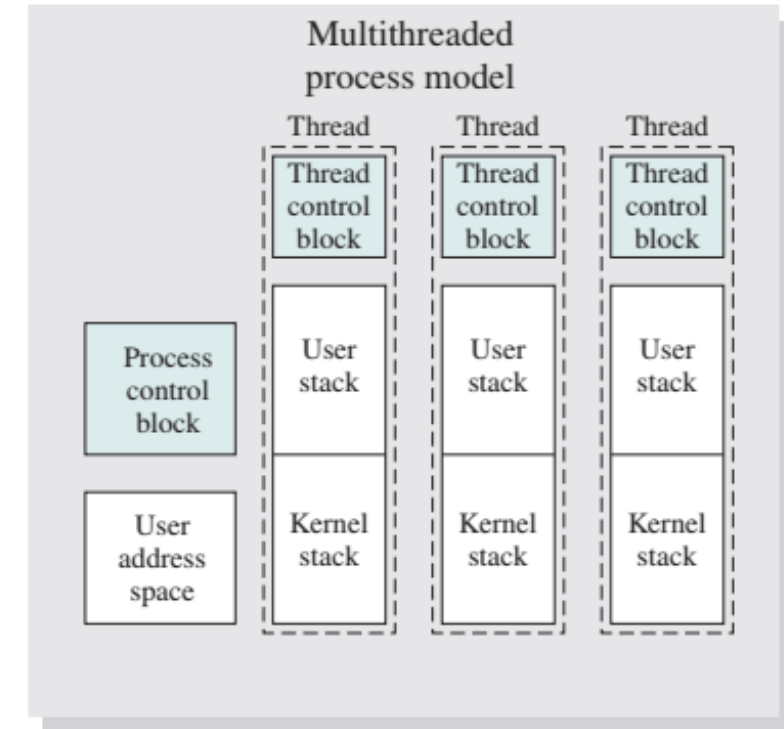
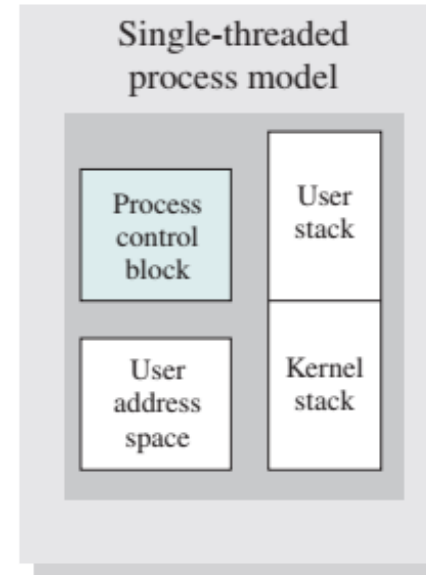


- Four fork() system calls; one for each process
- Each process executes independently
- IPC mechanism to communicate between processes
- Each process has its own instruction, data, heap, and stack



# Benefits of Threads

- Far less time to create a new thread in an existing process, than to create a brand-new process
- Less time to terminate a thread than a process
- Less time to switch between two threads within the same process than to switch between processes
- Threads enhance efficiency in communication between different executing programs



# Threads vs Processes

- A thread has no data segment or heap
- A thread cannot live on its own, it must live within a process
- There can be more than one thread in a process, the first thread calls `main()` & has the process's stack
- Inexpensive creation
- Inexpensive context switching
- Efficient communication
- If a thread dies, its stack is reclaimed
- A process has code/data/heap & other segments
- A process has at least one thread
- Threads within a process share code/data/heap, share I/O, but each has its own stack & registers
- Expensive creation
- Expensive context switching
- Interprocess communication can be expressive
- If a process dies, its resources are reclaimed & all threads die

Source: <http://www.cs.columbia.edu/~junfeng/13fa-w4118/lectures/l08-thread.pdf>


# Types of Threads

- **User-Level Threads**

- thread management is done by a user level thread library
- the kernel does not know anything about the threads running


- **Kernel-Level Threads**

- threads are directly supported by the kernels
- also known as light weight processes



# User Level Threads


- Fast as no system call to manage. Thread library does everything
- Switching is fast. NO switch from user to protected mode
- Scheduling can be an issue
- Lack of coordination between kernel and threads
- If one thread invokes a system call, all threads need to wait





# Kernel Level Threads

- Scheduler can decide to give more time to a process that large number of threads
- Since threads managed by kernel, no blocking on system calls
- Slow in comparison
- Overheads – scheduling threads apart from processes



# References

- William Stallings, “Operating Systems: Internals and Design Principles”, 9<sup>th</sup> edition, Pearson Edu. Ltd., 2018
- Charles Crowley, “Operating Systems: A design-oriented approach”, TMH
- Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau (University of Wisconsin-Madison), “Operating Systems: Three Easy Pieces”.  
URL: <http://pages.cs.wisc.edu/~remzi/OSTEP/>

# Process Synchronization Deadlocks Inter Process Communication (IPC)

CS3003D: Operating Systems

# Concurrency

- Concurrency encompasses a host of design issues, including
  - communication among processes
  - sharing of and competing for resources (such as memory, files, and I/O access)
  - synchronization of the activities of multiple processes, and allocation of processor time to processes
- Concurrency arises in three different contexts
  - Multiple applications
  - Structured applications
  - Operating system structure

# Synchronization

shared variable  
int flag = 5

ProgramA

flag++

Output value of flag can be 5, 4, or 6 based on the way the processes are executing, when context switching happens

ProgramB

flag --

- 1) reg1 = flag
- 2) reg1 = reg1 + 1
- 3) flag = reg1

- 4) reg2 = flag
- 5) reg2 = reg2 - 1
- 6) flag = reg2

## Scenario1

ProcessA

- 1)
  - 2)
  - 3) flag = 6
- Context Switch

ProcessB

- 4)
- 5)
- 6) flag = 5

## Scenario2

ProcessB

- 1)
  - 2)
  - 3) flag = 4
- Context Switch

ProcessA

- 4)
- 5)
- 6) flag = 5

## Scenario3

ProcessA

- 1) reg1=5
- Context Switch

ProcessB

- 2) reg2 = 5
  - 3) reg2 = 4
  - 4) flag = 4
- Context Switch

ProcessA

- 5) reg1 = 6
- 6) flag = 6

## Scenario4

ProcessB

- 1) reg2=5
- Context Switch

ProcessA

- 2) reg2 = 5
  - 3) reg2 = 6
  - 4) flag = 6
- Context Switch

ProcessB

- 5) reg1 = 4
- 6) flag = 4

```

1 #include <stdio.h>
2 #include <pthread.h>
3 #include "common.h"
4 #include "common_threads.h"
5
6 static volatile int counter = 0;
7
8 // mythread()
9 //
10 // Simply adds 1 to counter repeatedly, in a loop
11 // No, this is not how you would add 10,000,000 to
12 // a counter, but it shows the problem nicely.
13 //
14 void *mythread(void *arg) {
15     printf("%s: begin\n", (char *) arg);
16     int i;
17     for (i = 0; i < 1e7; i++) {
18         counter = counter + 1;
19     }
20     printf("%s: done\n", (char *) arg);
21     return NULL;
22 }
23
24 // main()
25 //
26 // Just launches two threads (pthread_create)
27 // and then waits for them (pthread_join)
28 //
29 int main(int argc, char *argv[]) {
30     pthread_t p1, p2;
31     printf("main: begin (counter = %d)\n", counter);
32     Pthread_create(&p1, NULL, mythread, "A");
33     Pthread_create(&p2, NULL, mythread, "B");
34
35     // join waits for the threads to finish
36     Pthread_join(p1, NULL);
37     Pthread_join(p2, NULL);
38     printf("main: done with both (counter = %d)\n",
39           counter);
40     return 0;
41 }

```

```

prompt> gcc -o main main.c -Wall -pthread; ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 20000000)

```

```

prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19345221)

```

```

prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19221041)

```

Source: Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau (University of Wisconsin-Madison),  
*"Operating Systems: Three Easy Pieces"*.

# Race condition

- Many processes manipulate the same data portion
- During concurrent execution, outcome depends upon the order in which the access happens
- Incorrect data leads to misleading output
- Can be prevented by synchronization between processes

How to avoid race condition?

- Prohibit more than one process from reading and writing the shared data (critical section) at the same time

# Definitions

- Critical Section
  - A section of code within a process that requires access to shared resources, and that must not be executed while another process is in a corresponding section of code
- Race condition
  - A situation in which multiple threads or processes read and write a shared data item, and the final result depends on the relative timing of their execution
- Mutual exclusion
  - The requirement that when one process is in a critical section that accesses shared resources, no other process may be in a critical section that accesses any of those shared resources



# Operating System concerns

- OS must be able to keep track of the various processes
  - PCB
- OS must allocate and deallocate resources for each active process
  - Processor timer, Memory, Files, I/O devices
- Protection of data and resources of each process against unintended interference by other processes

# Process Interaction

(Competition for resources and Cooperation among processes)

Degree of Awareness	Relationship	Influence that One Process Has on the Other	Potential Control Problems
Processes unaware of each other	Competition	<ul style="list-style-type: none"><li>• Results of one process independent of the action of others</li><li>• Timing of process may be affected</li></ul>	<ul style="list-style-type: none"><li>• Mutual exclusion</li><li>• Deadlock (renewable resource)</li><li>• Starvation</li></ul>
Processes indirectly aware of each other (e.g., shared object)	Cooperation by sharing	<ul style="list-style-type: none"><li>• Results of one process may depend on information obtained from others</li><li>• Timing of process may be affected</li></ul>	<ul style="list-style-type: none"><li>• Mutual exclusion</li><li>• Deadlock (renewable resource)</li><li>• Starvation</li><li>• Data coherence</li></ul>
Processes directly aware of each other (have communication primitives available to them)	Cooperation by communication	<ul style="list-style-type: none"><li>• Results of one process may depend on information obtained from others</li><li>• Timing of process may be affected</li></ul>	<ul style="list-style-type: none"><li>• Deadlock (consumable resource)</li><li>• Starvation</li></ul>

# Three requirements for critical section problem

- Mutual Exclusion

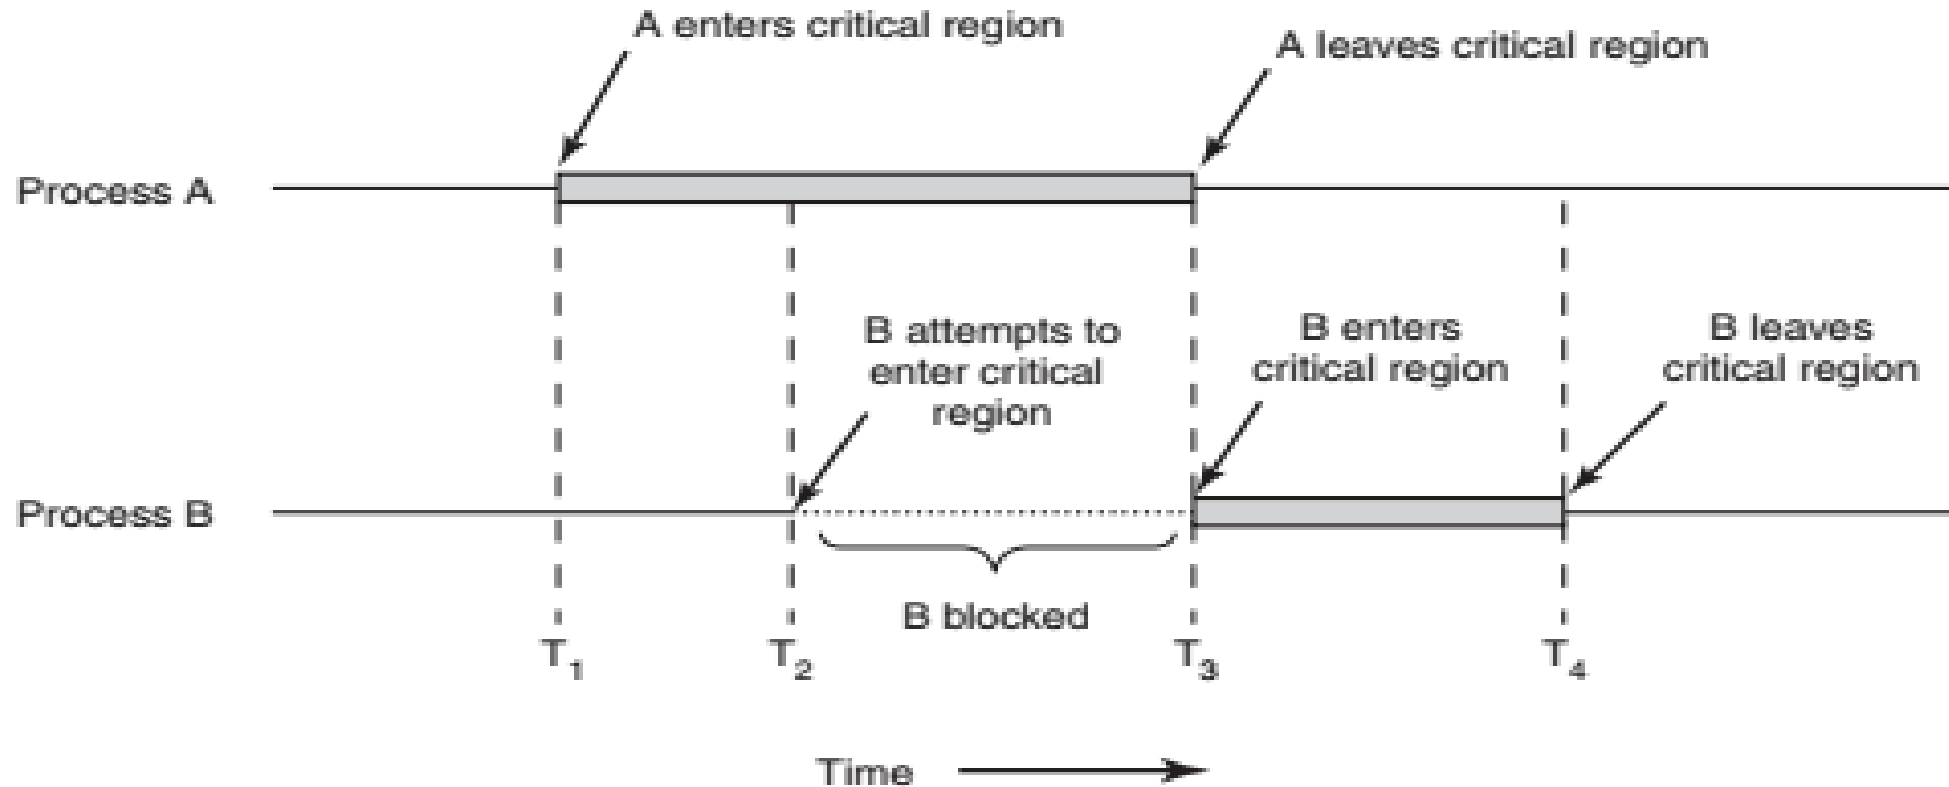
- No two processes may be simultaneously inside their critical regions

- Progress

- No process running outside its critical region may block any process

- No starvation (bounded waiting)

- No process should have to wait forever to enter its critical region



# Solutions to critical section

- Disable interrupts

- Context switches will not happen
- Codes that execute in the kernel can only disable interrupts
- User processes/application programs cannot disable interrupts

```
While(TRUE) {  
    // code area  
    disable_interrupts() < LOCK  
    critical_section  
    enable_interrupts() < UNLOCK  
    // other code area  
}
```

# Busy waiting

Process-1

```
while (TRUE) {  
    while (turn == 2);    // LOCK  
    critical_section  
    turn = 2;            // UNLOCK  
    reminder_code_here  
}
```

Shared  
int turn = 1;

Process-2

```
while (TRUE) {  
    while (turn == 1);    // LOCK  
    critical_section  
    turn = 1;            // UNLOCK  
    reminder_code_here  
}
```

- Mutual exclusion achieved
- Busy waiting – resource wastage
  - When Process-2 executes first, always in loop; always in primary memory – either at READY state or at RUNNING state
- Progress condition is violated

# No Mutual Exclusion

Shared

p1\_inside = false, p2\_inside = false

Process-1

```
while (TRUE) {  
    while (p2_inside == TRUE);    // LOCK  
    p1_inside = TRUE;  
    critical_section  
    p1_inside = FALSE;            // UNLOCK  
}
```

Process-2

```
while (TRUE) {  
    while (p1_inside == TRUE);    // LOCK  
    p2_inside = TRUE;  
    critical_section  
    p2_inside = FALSE;            // UNLOCK  
}
```

```
while(p2_inside == TRUE);  
// Context Switch (Process-2)  
while(p1_inside == TRUE);  
p2_inside = TRUE;  
// Context Switch (Process-1)  
p1_inside = TRUE;
```

- Mutual exclusion is not guaranteed
- Both processes can enter into critical section

# Endless Wait

Shared

p1\_inside = false, p2\_inside = false

Process-1

Process-2

```
while (TRUE) {  
    p1_inside = TRUE;  
    while (p2_inside == TRUE);    // LOCK  
    critical_section  
    p1_inside = FALSE;            // UNLOCK  
}
```

```
while (TRUE) {  
    p2_inside = TRUE;  
    while (p1_inside == TRUE);    // LOCK  
    critical_section  
    p2_inside = FALSE;            // UNLOCK  
}
```

```
p1_inside = TRUE  
// Context Switch (Process-2)  
p2_inside = TRUE;
```

- Achieves Mutual exclusion
- Can it progress?
  - DEADLOCK!

# Peterson's Solution

Process-1

Shared

p1\_inside, p2\_inside, favoured

Process-2

```
while (TRUE) {  
    p1_inside = TRUE;  
    favoured = 2;  
    while (p2_inside == TRUE AND favoured = 2); // LOCK  
    critical_section  
    p1_inside = FALSE;           // UNLOCK  
}
```

```
while (TRUE) {  
    p2_inside = TRUE;  
    favoured = 1;  
    while (p1_inside == TRUE AND favoured = 1); //  
    LOCK  
    critical_section  
    p2_inside = FALSE;           // UNLOCK  
}
```

- Breaking the deadlock as one of the processes is favoured
- Solves critical section problem for two processes.



# Bakery Algorithm – synchronization between $N$ processes ( $N > 2$ )

- Proposed by Leslie Lamport
- Similar to token system in the bakeries/banks
  - Customers upon entering the bank is issued with the token
  - Waits until his/her turn arrives
  - Dispense the token and the service is rendered

Ref.: <http://www.cs.umd.edu/~shankar/412-S99/note-7.html>

# Simplified version of Bakery algorithm

- Each process is numbered 0 to N-1
- Each process i has an integer variable num[i], initially 0, that is readable by all processes but writeable by process i only

```
Entry(i) {  
    num[i] = MAX( num[0], num[1], ... , num[N-1] ) + 1 ;  
    for p = 0 to N-1 do {  
        while (num[p] != 0 AND num[p] < num[i]) do no-op ;  
    }  
}
```

Lock

Critical section

```
Exit(i) {  
    num[i] = 0 ;  
}
```

Unlock

# Example

num[i]	P0	P1	P2	P3	P4
Initial	0	0	0	0	0
P2	0	0	1	0	0
P3	0	0	0	2	0
P4	0	0	0	0	3
P0	4	0	0	0	0
P1	4	5	0	0	0
Final	4	5	1	2	3

$$\text{num}[i] = \text{MAX}(\text{num}[0], \text{num}[1], \dots, \text{num}[N-1]) + 1$$

num[i]	P0	P1	P2	P3	P4
Initial	4	5	1	2	3
P2	4	5	0	2	3
P3	4	5	0	0	3
P4	4	5	0	0	0
P0	0	5	0	0	0
P1	0	0	0	0	0
Final	0	0	0	0	0

```
for p = 0 to N-1 do {  
    while(num[p] != 0 AND num[p] < num[i]) do no-op ;  
}
```

## Problem!

**Assumption: No two processes get the same token**

When two process gets the same num[i] value (same token)

Two processes enter into the critical section